e-Elves

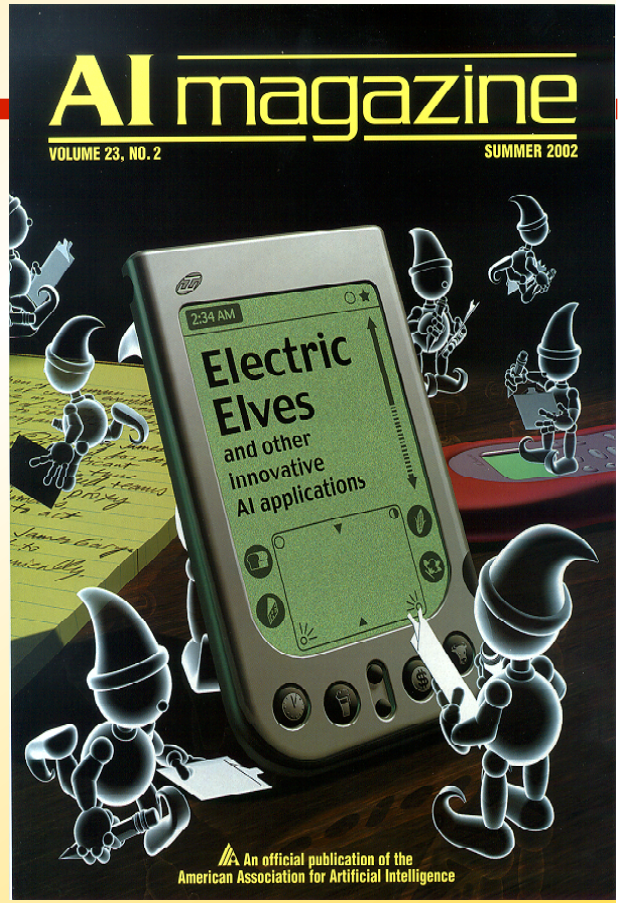

Beyond the Elves: Making Intelligent Agents Intelligent

Craig Knoblock

University of Southern California
Information Sciences Institute

Contributors

- José Luis Ambite, USC
- Greg Barish, Fetch Technologies
- Hans Chalupsky, USC
- Yolanda Gil, USC
- Kristina Lerman, USC
- Steven Minton, Fetch Technologies
- Maria Muslea, USC
- Jean Oh, CMU
- David V. Pynadath, USC
- Thomas A. Russ, USC
- Paul Scerri, USC
- Milind Tambe, USC

Research Support

- Electric Elves
  - DARPA CoABS program
  - DARPA Active Templates program
- Research based on lessons learned
  - DARPA PAL Program
  - AFOSR

Outline

- The Electric Elves
  - Overview
- The Travel Elves
  - Constraint-based data integration
  - Efficient plan execution
- Beyond the Elves
  - Deployment to DARPA
  - Some lessons
    - Learning expectations
    - Finding alternative sources
    - Predicting prices
- Discussion

Agent-Facilitated Human Organizations

Goal

Agent-Facilitated Human Organizations

- Develop software agents that automate routine tasks within organization
  - E.g., obtaining info, monitoring, distributing information
- Enable software agents and humans to act coherently within the organization
  - Humans have agent proxies that assist in coordinating with other agents
- Coordination of tasks within the organization
  - Efficient use of resources (physical and human)
  - Monitoring progress of both individual and overall tasks
  - Suggest/execute corrective actions when goals appear endangered

Example Tasks

Supporting our own research organization:

- Coordinating routine project activities
  - Eliminate email about meeting delays/postponements
  - Eliminate email about who will present at next meeting
  - Track visitor's flights and arrival times
- Support for external meetings
  - Monitor all aspects of a traveler's plans
    - Notify traveler of prices changes, schedule changes, flight delays
    - Send fax to hotel if traveler is delayed
  - Organize meetings for people with similar research interests in other organizations
  - Organize meeting at external site (e.g., a restaurant)
- Similar issues arise in many organizations: corporations, military, etc.

Agent Organization is Grounded in the Real World

- On-line information sources
  - Calendar and schedule information
  - Flight schedules, restaurants, etc.
- "Sensors" and "actuators"
  - GPS, "finger", infrared, etc.
  - Postpone/cancel meetings, faxes to hotels and restaurants, etc.
- Interaction with people
  - Communicate through personal portable devices (PDAs, phones, etc.)

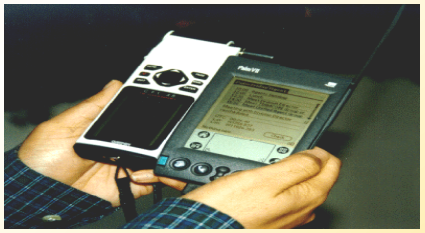

Summary

- Electric Elves: successful use of agent technology to support human organization
  - Interact with their environment:
    - Data: on-line information sources, calendars, GPS, etc.
    - Interaction with people: faxes, email, PDAs, phones, etc.
  - Support routine activities
    - Arrange meetings, lunch orders
    - Monitoring and suggesting/executing corrective actions
- Integrates powerful techniques in diverse areas within AI:
  - Teamwork, human proxies [Teamcore]
  - Ontologies & problem solving [Expect]
  - Information agents [Ariadne]
  - Knowledge representation & information retrieval [PowerLoom]

Outline

- The Electric Elves
  - Overview
- The Travel Elves
  - Constraint-based data integration
  - Efficient plan execution
- Beyond the Elves
  - Deployment to DARPA
  - Some lessons
    - Learning expectations
    - Finding alternative sources
    - Predicting prices
- Discussion

Interactive Trip Planning

- Commercial systems provide support to select flights, hotels and cars
  - Integrates the planning at the level of dates and locations
- There are many more factors involved in planning a trip
  - Which airports to fly into and out of
  - Whether to drive or take a taxi to the airport
  - How to get from the airport to the destination
  - Proximity of hotel to meeting
  - Etc…
- Ideally a system will
  - Provide all of the data required to make these decisions
  - Provide a way to consider the tradeoffs of the various choices

Heracles Constraint-based Planning

- Framework for building integrated applications
- Extract and integrate data for a given task
  - Live access to online sources using the wrappers
- Constraint-based system decides what sources to query and how to integrate the results
  - Tight integration of user choices

The Travel Assistant

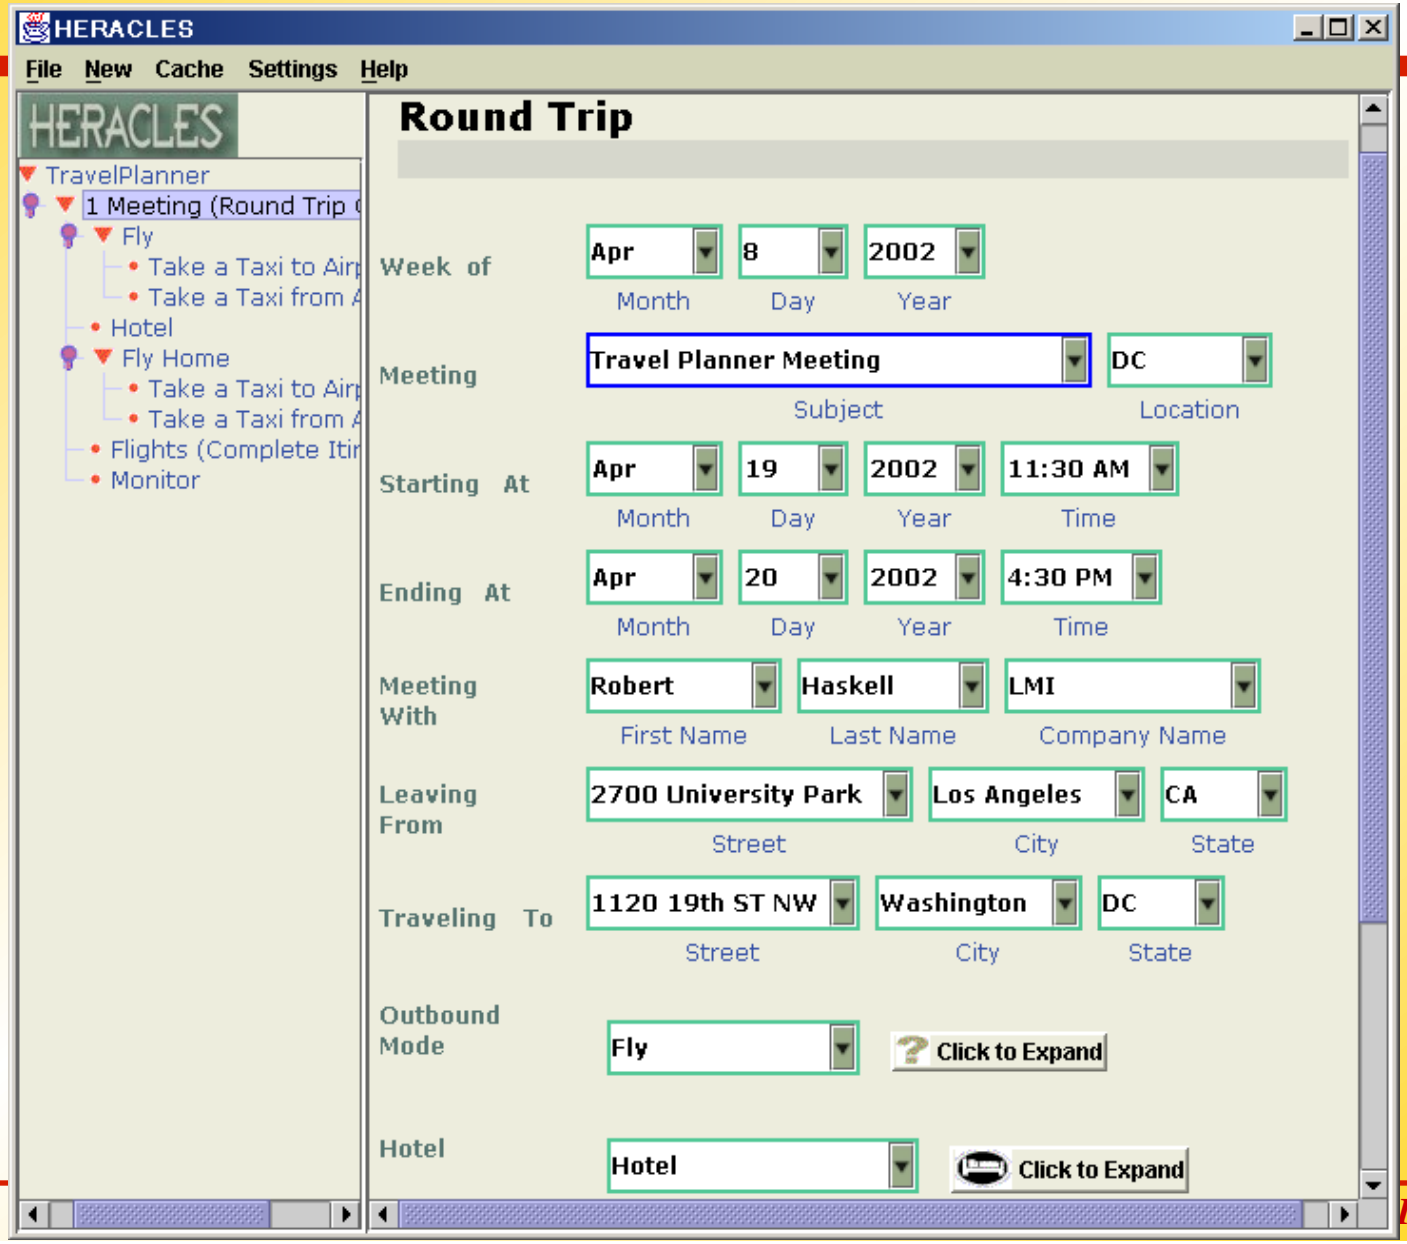

Supports Informed Choices

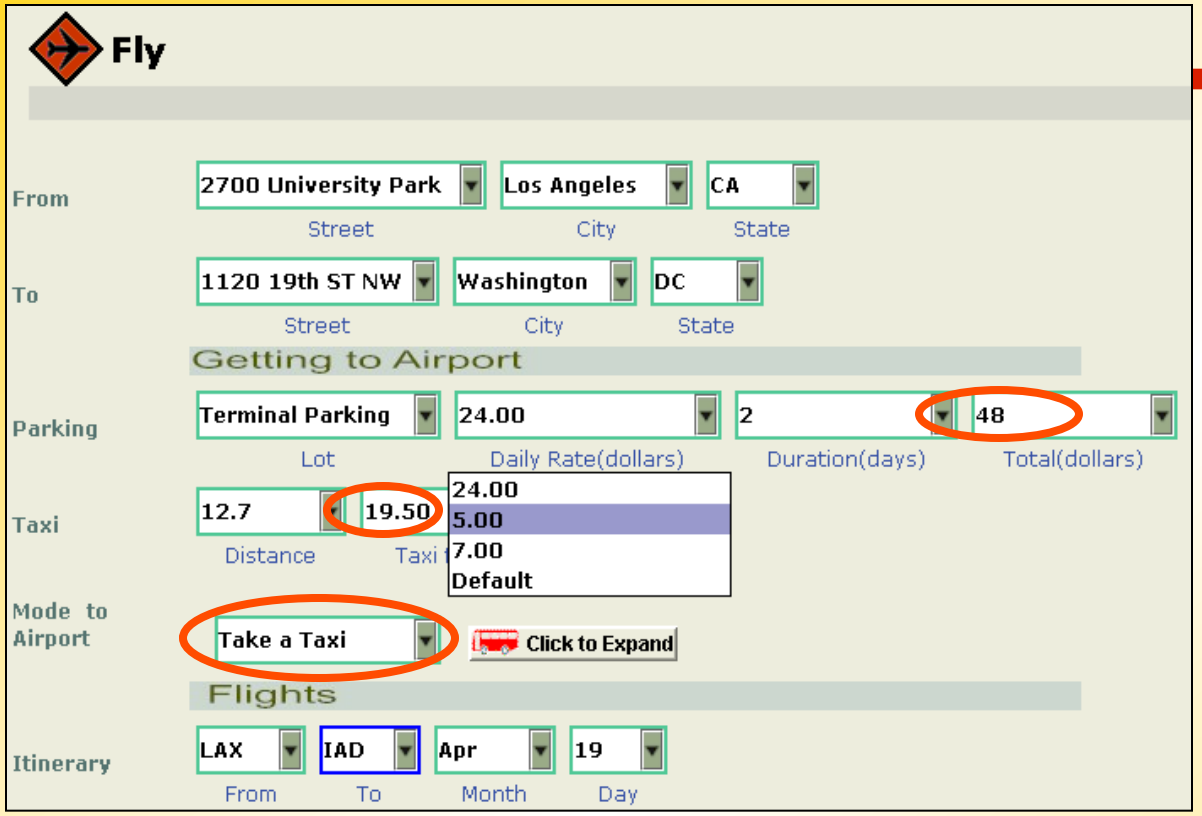

Supports Informed Choices

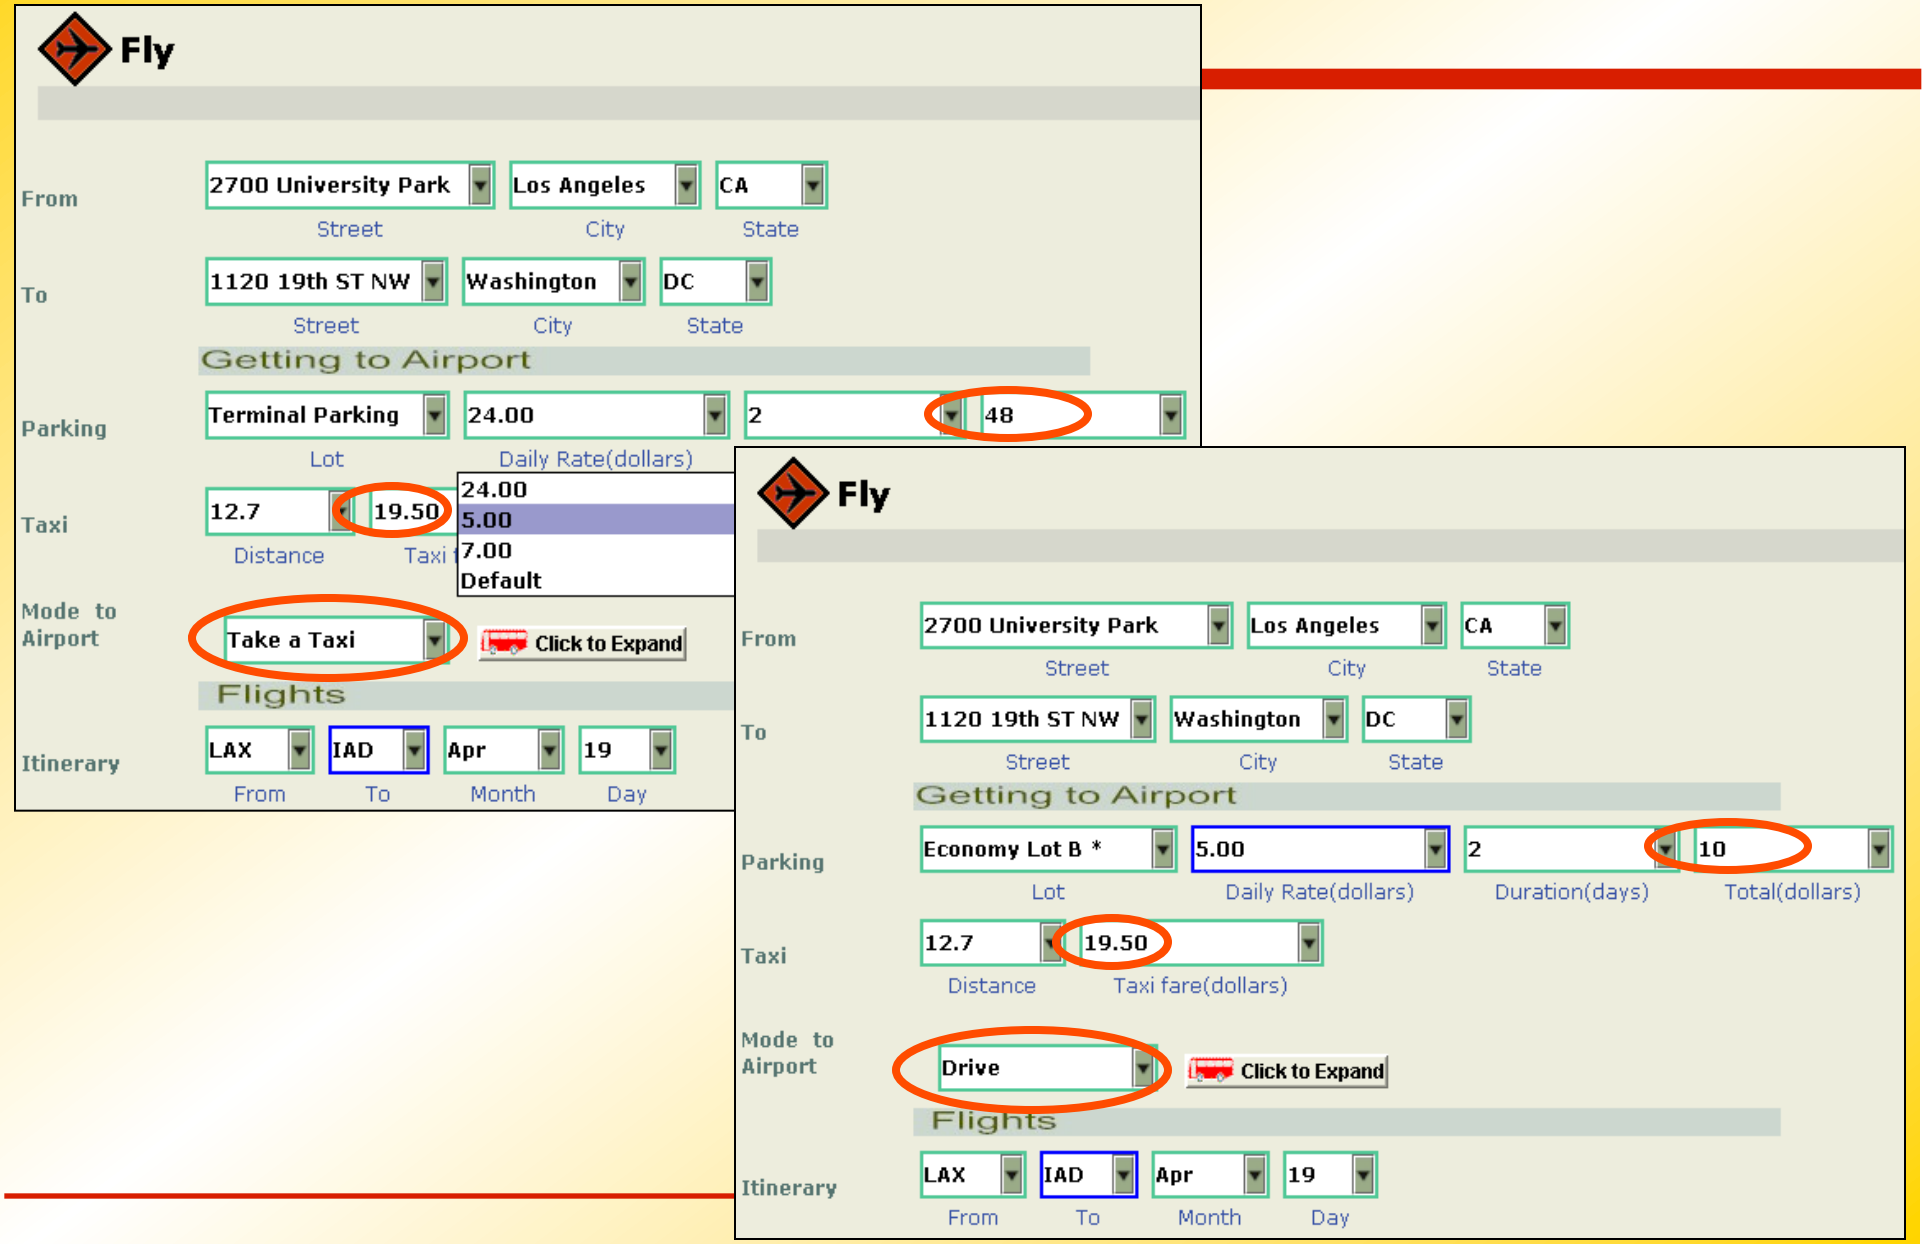

Changes Propagate Throughout

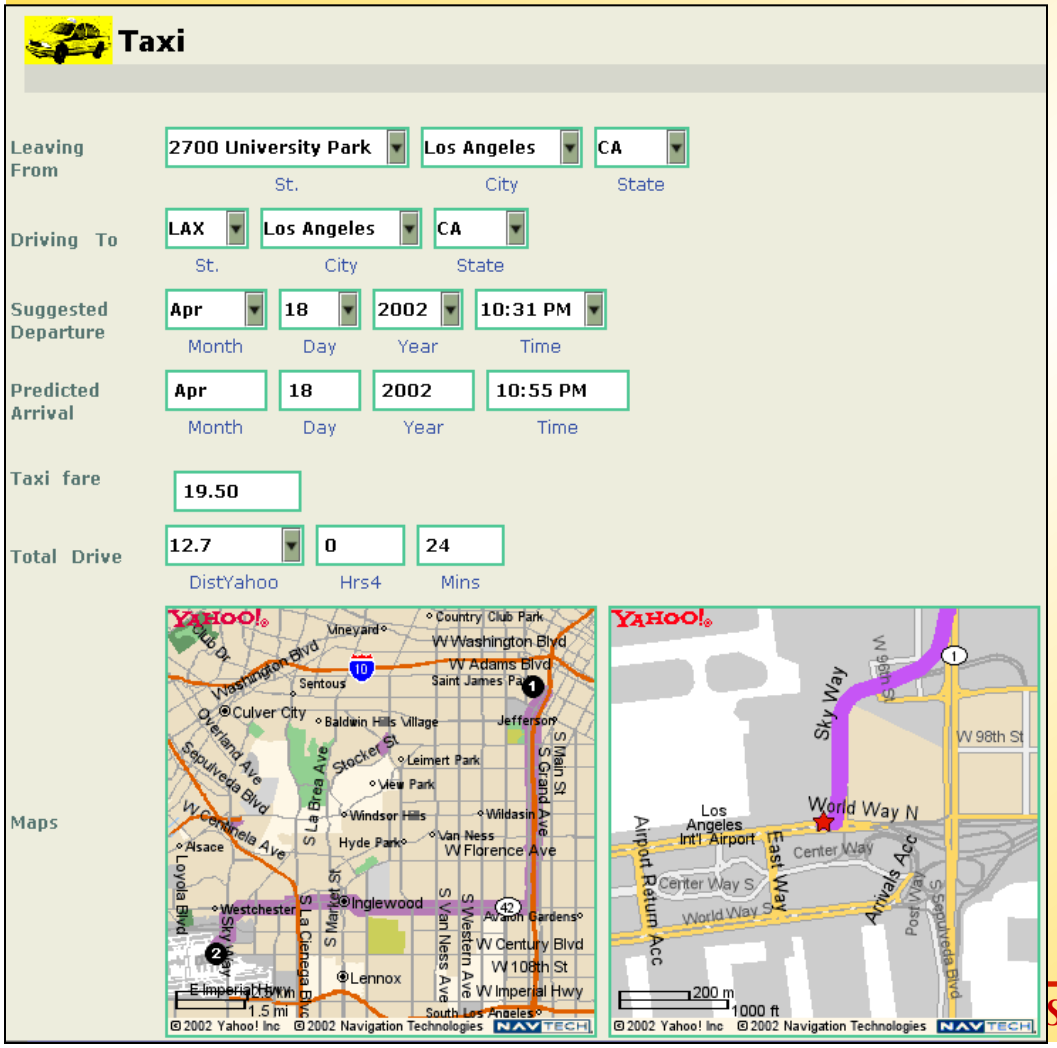

Changes Propagate Throughout

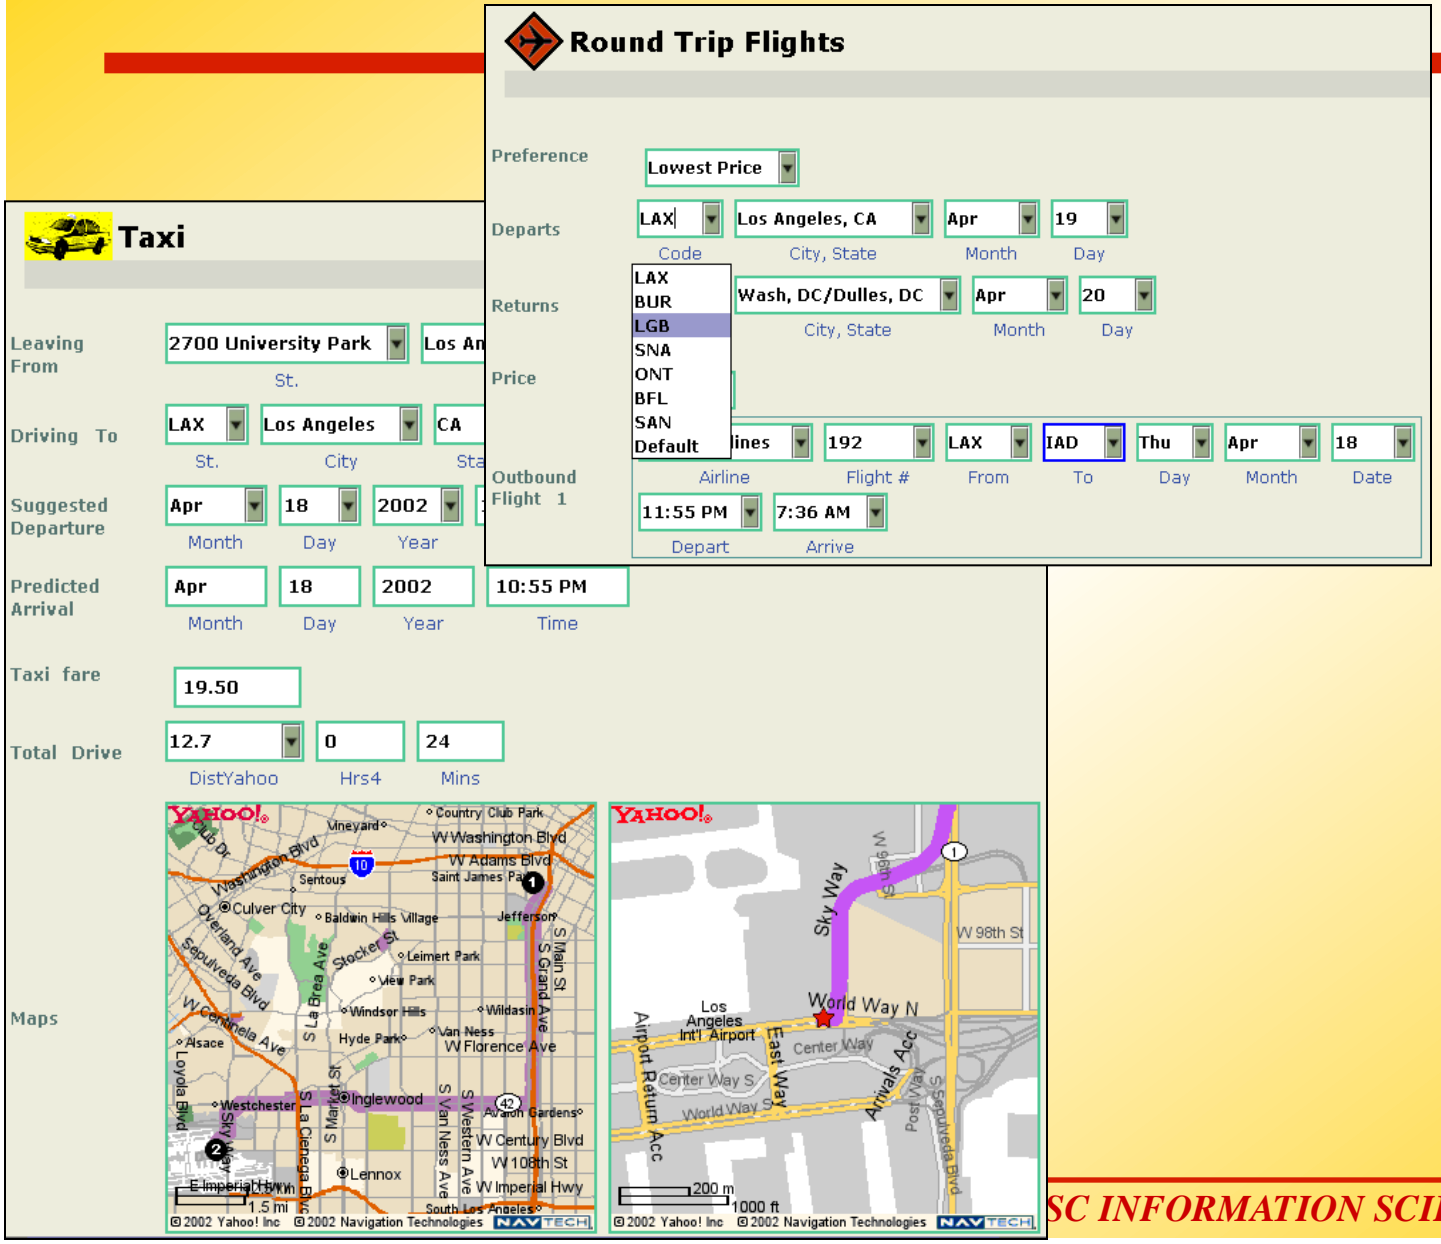

Changes Propagate Throughout

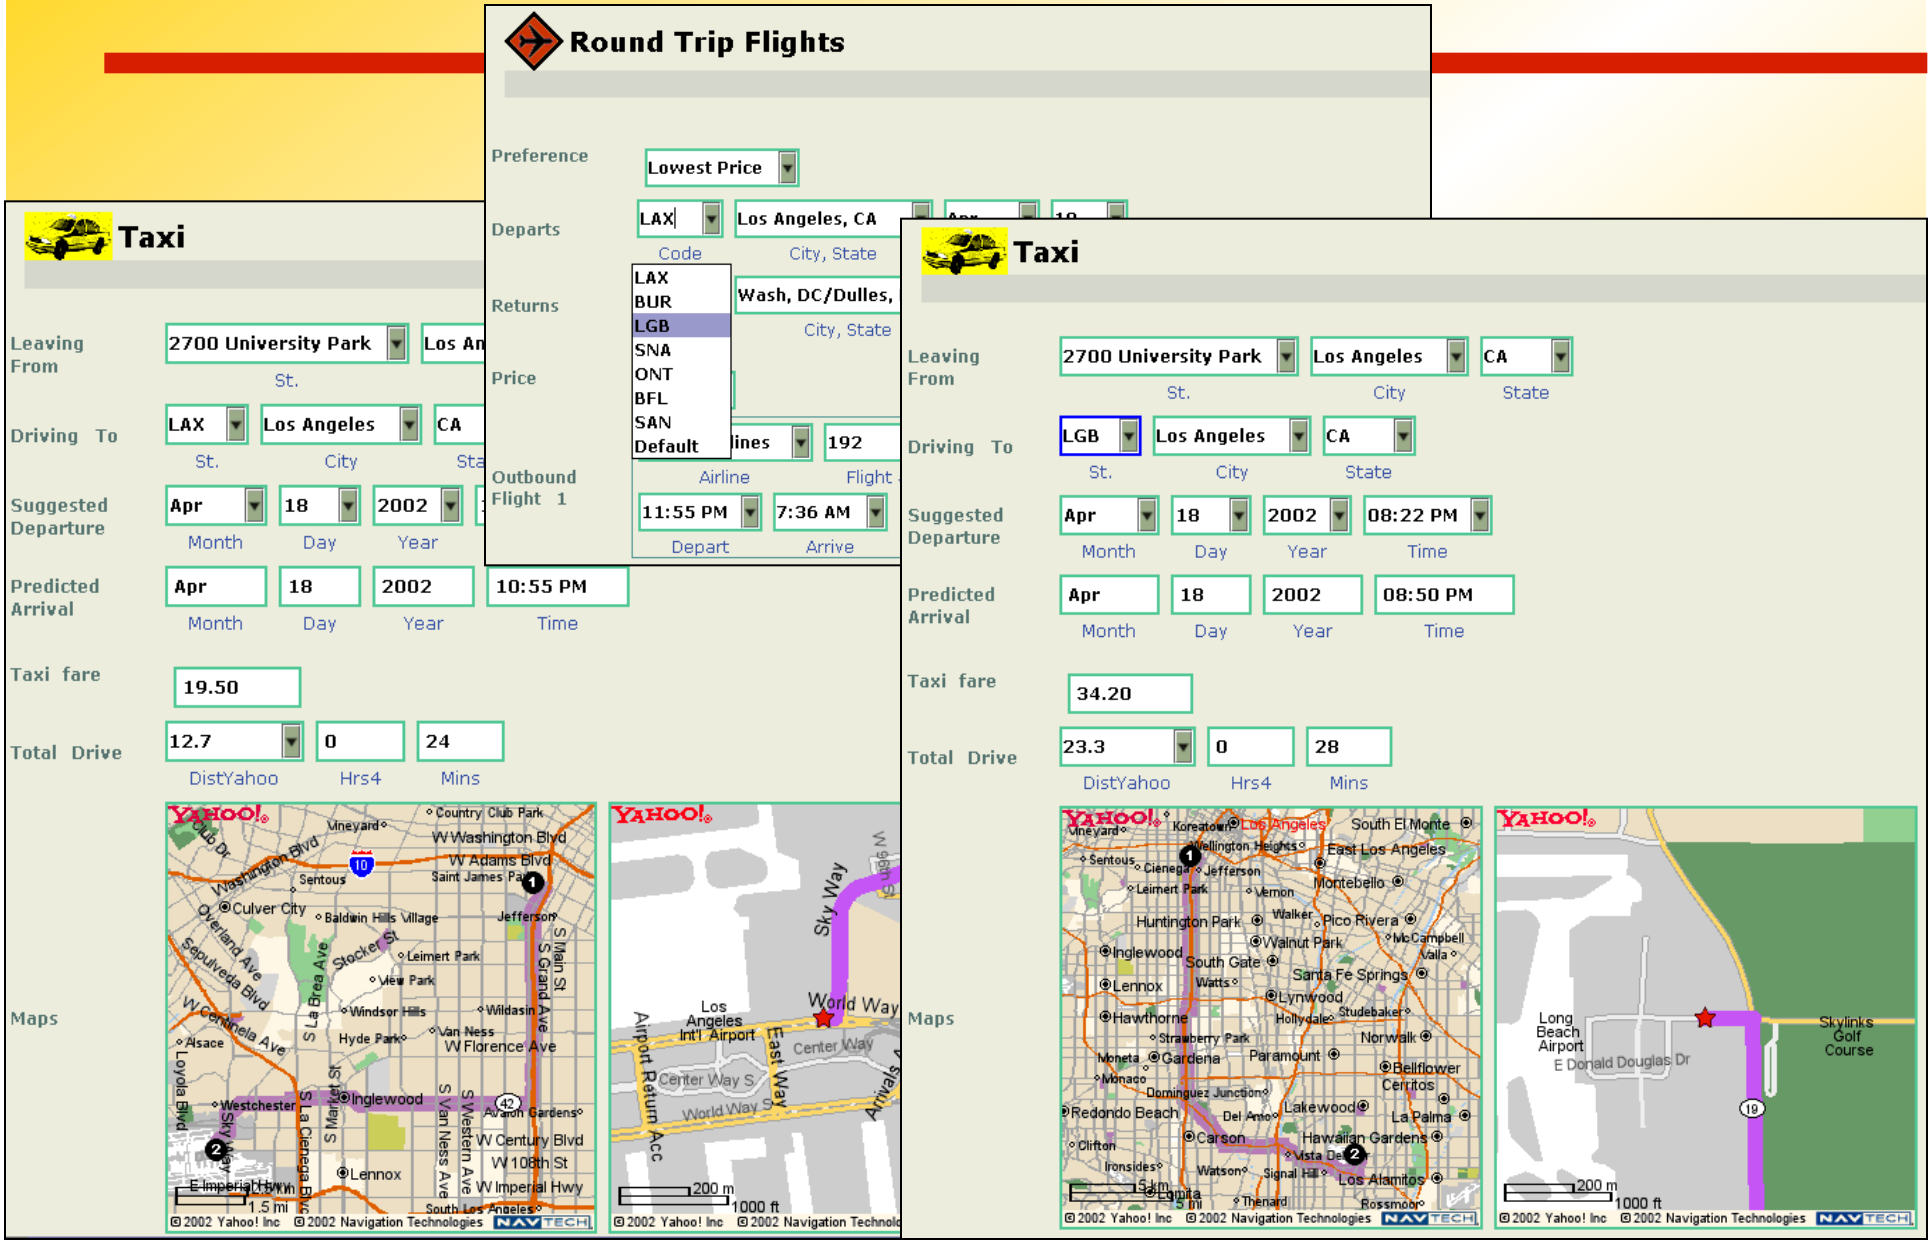

Constraint Network: Drive or Taxi?

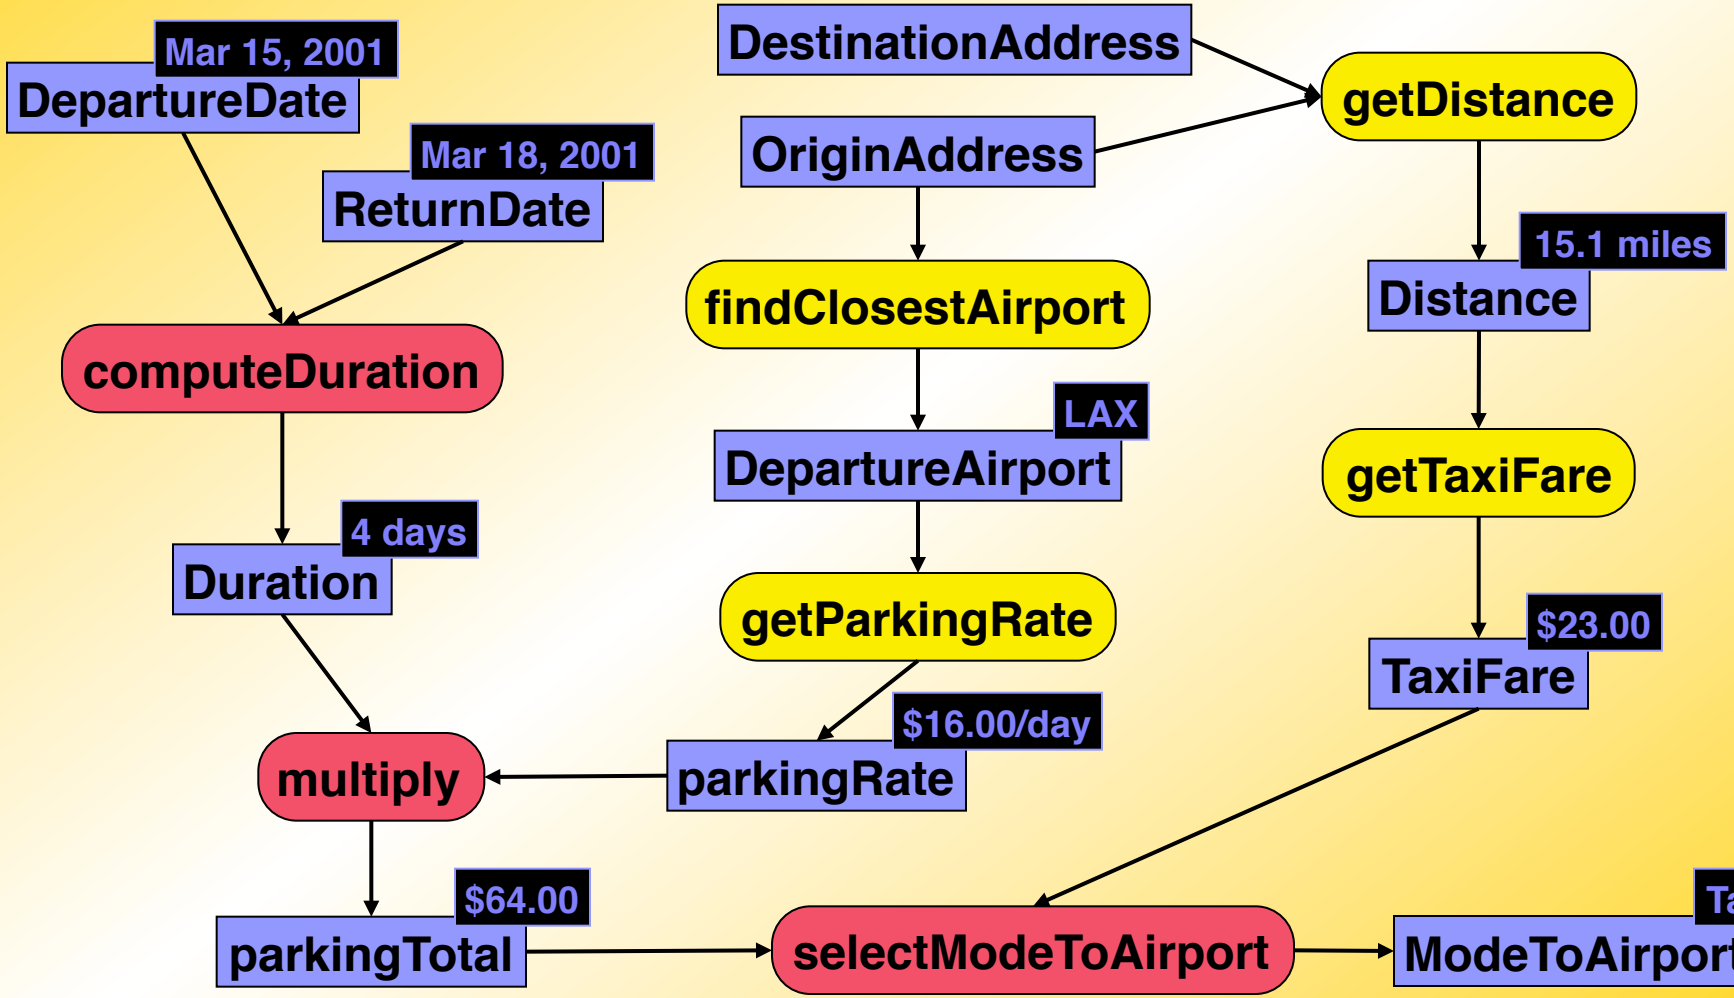

Summary

- Integration of wide range of data from many different sources
- Tight integration of data using constraints to capture the dependencies
- Supports better decision making
  - Easy to consider costs of specific choices
  - Easy to compare tradeoffs

Outline

- The Electric Elves
  - Overview
- The Travel Elves
  - Constraint-based data integration
  - Efficient plan execution
- Beyond the Elves
  - Deployment to DARPA
  - Some lessons
    - Learning expectations
    - Finding alternative sources
    - Predicting prices
- Discussion

Agents for Monitoring Travel

- Many opportunities and possible problems can arise during travel
- Current environment:
  - Wide access to data
  - Abundance of computer resources
  - Availability of cell phones and portable computers
- Makes it possible to monitor all aspects of a trip
- Create personal assistants that monitor your travel plan to
  - exploit opportunities
  - avoid problems

Monitoring Travel Plans

Monitoring Tasks

Monitor Flight Status
- (•) Monitor Flights
- ( ) Stop Monitoring

7038128516 — Notify Hotel (Fax)

7034948462 — Notify Car Rental Counter (Fax)

Status

Outbound flight 1	Outbound flight 2	Inbound flight 1	Inbound flight 2
Active	Active	Active	Active

Monitor Flight Schedule
- (•) Monitor Schedule
- ( ) Stop Monitoring

Active — Status

Monitor Earlier Flights
- (•) Monitor Earlier Flights
- ( ) Stop Monitoring

Active — Status

Monitor Connecting Flights
- (•) Monitor Connecting Flights
- ( ) Stop Monitoring

Active — Status (Outbound)

Active — Status (Inbound)

Monitor Airfare

Decrease only — Mode

- (•) Monitor Airfare
- ( ) Stop Monitoring

Airfare

Active — Status

Agents Deployed to Monitor Travel Itinerary

Travel Itinerary

GRID

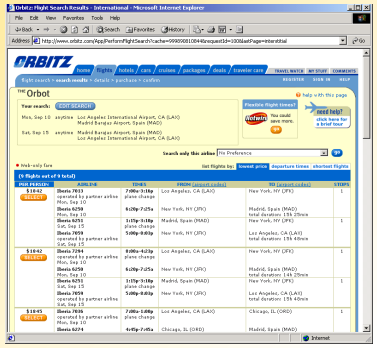

Flight Prices & Schedules

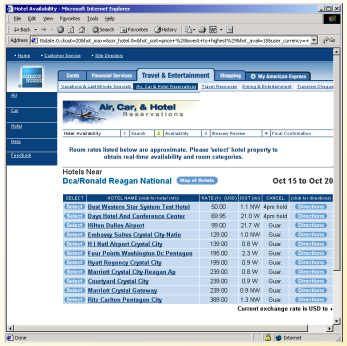

Flight Status

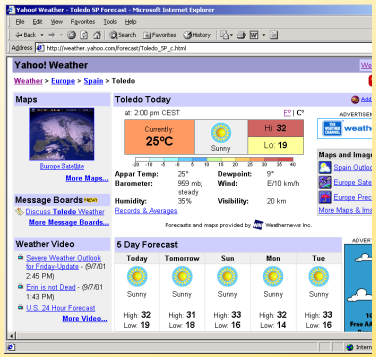

Weather

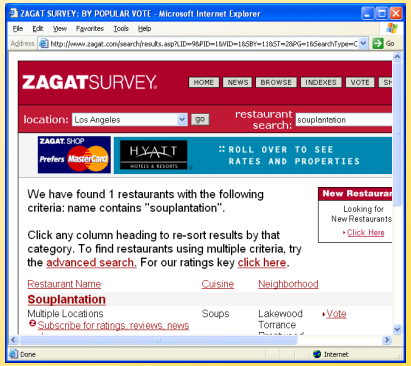

Restaurants

Example Agents

- Flight-Status Agent:
  - Flight delayed message:

```
Your United Airlines flight 190 has been delayed.
It was  originally scheduled to depart at 11:45 AM and
is now scheduled to depart at 12:30 PM.
The new arrival time is 7:59 PM.
```

  - Flight cancelled message:

```
Your Delta Air Lines flight 200 has been cancelled.
```

  - Fax to hotel message:

```
Attention: Registration Desk
I am sending this message on behalf of David Pynadath,
who has a reservation at your hotel. David Pynadath is
on United Airlines 190, which is now scheduled to
arrive at IAD at 7:59 PM. Since the flight will be
arriving late, I would like to request that you
indicate this in the reservation so that the room is
not given away.
```

Monitoring Agents

- Airfare Agent: Airfare dropped message

```
The airfare for your American Airlines itinerary
(IAD - LAX) dropped to $281.
```

- Earlier-Flight Agent: Earlier flights message

```
The status of your currently scheduled flight is:
# 190 LAX (11:45 AM) - IAD (7:29 PM) 45 minutes Late
If you would like to return earlier, the following
United Airlines flights will arrive earlier than your
scheduled flights:
# 946 LAX (8:31 AM) - IAD (3:35 PM) 11 minutes Late
--------
# 388 LAX (9:25 AM) - DEN (12:25 PM) 10 minutes Late
# 1534 DEN (1:20 PM) - IAD (6:06 PM) On Time
```

Theseus

- A plan language and execution system for building Web-based information-gathering and monitoring agents [Barish & Knobock, JAIR'05]
  - *Efficient* enough for near-real-time monitoring: streaming dataflow
  - *Expressive* enough for integrating a variety of sources (web sites, XML, databases, …)

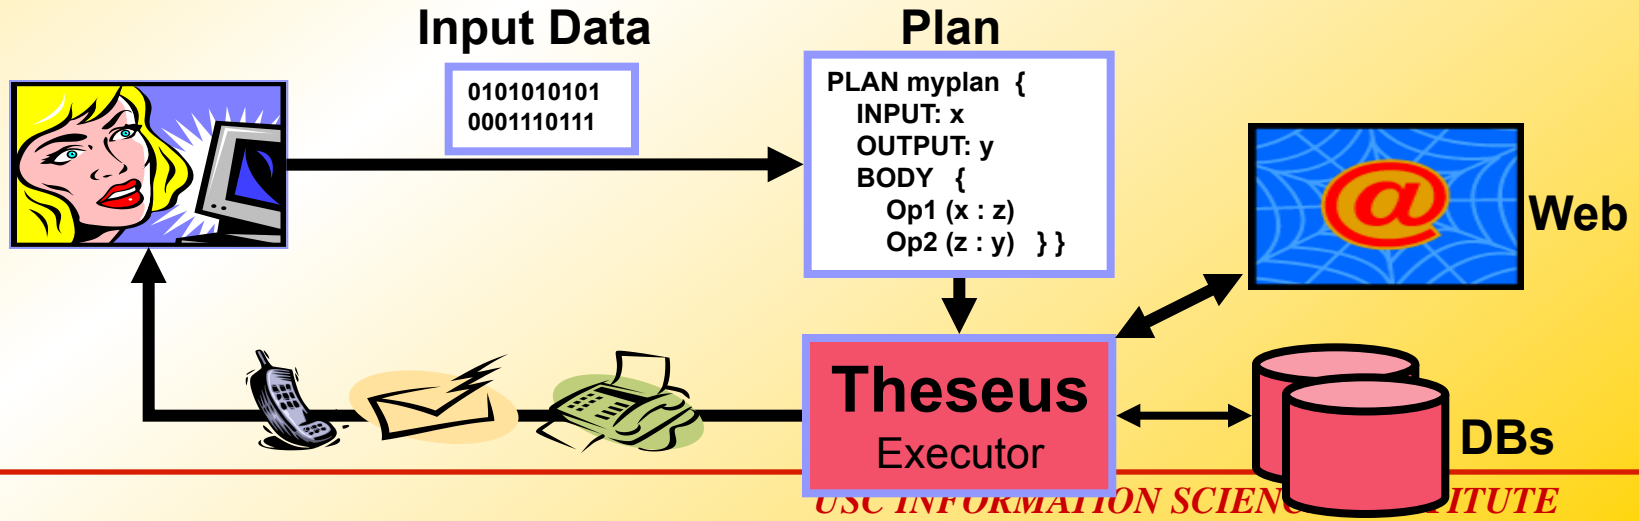

Streaming Dataflow

- Plans consist of a network of operators
  - Examples: **Wrapper**, **Select**, etc.
  - Operators produce and consume data
  - Operators "fire" upon any input data

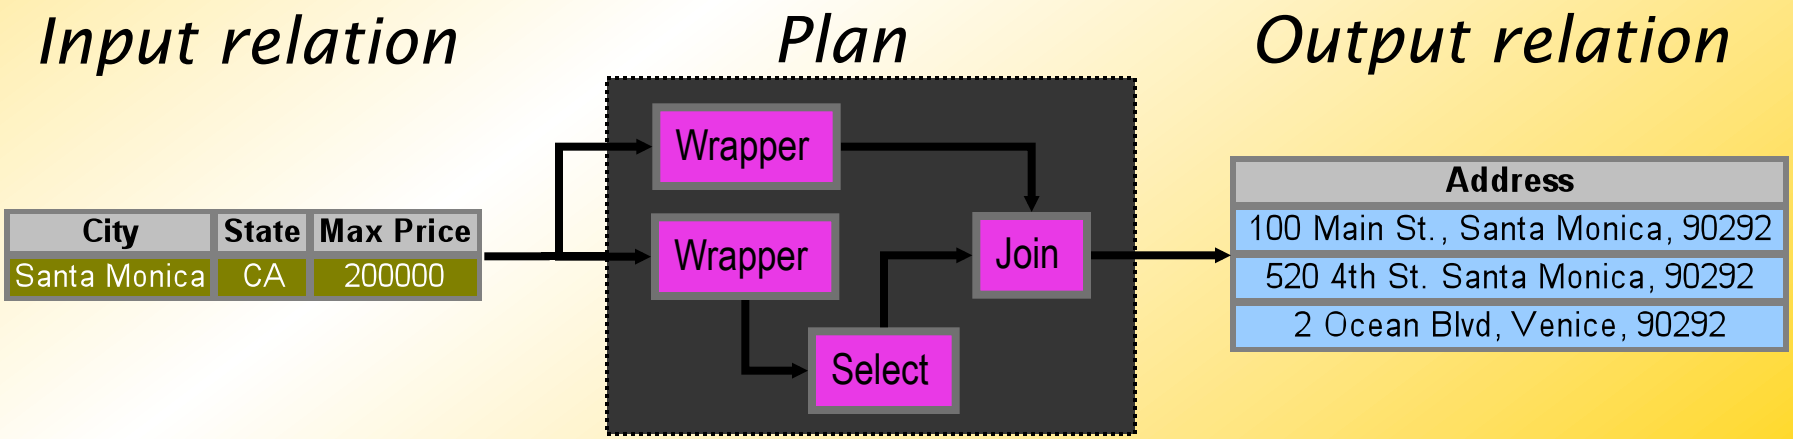

Theseus: Efficiency Streaming Dataflow

- Dataflow-style execution
  - Operators execute when inputs become available
  - Optimizes *horizontal parallelism*
    - Plan is as parallel as its data dependencies allow
- Data Streaming
  - Data in the system represented as *relations*
    - Producer operators pipeline *tuples* to consumers
  - Optimizes *vertical parallelism*
    - Multiple operators can work on same relation concurrently

Theseus Monitoring Agent: Flight Status Plan

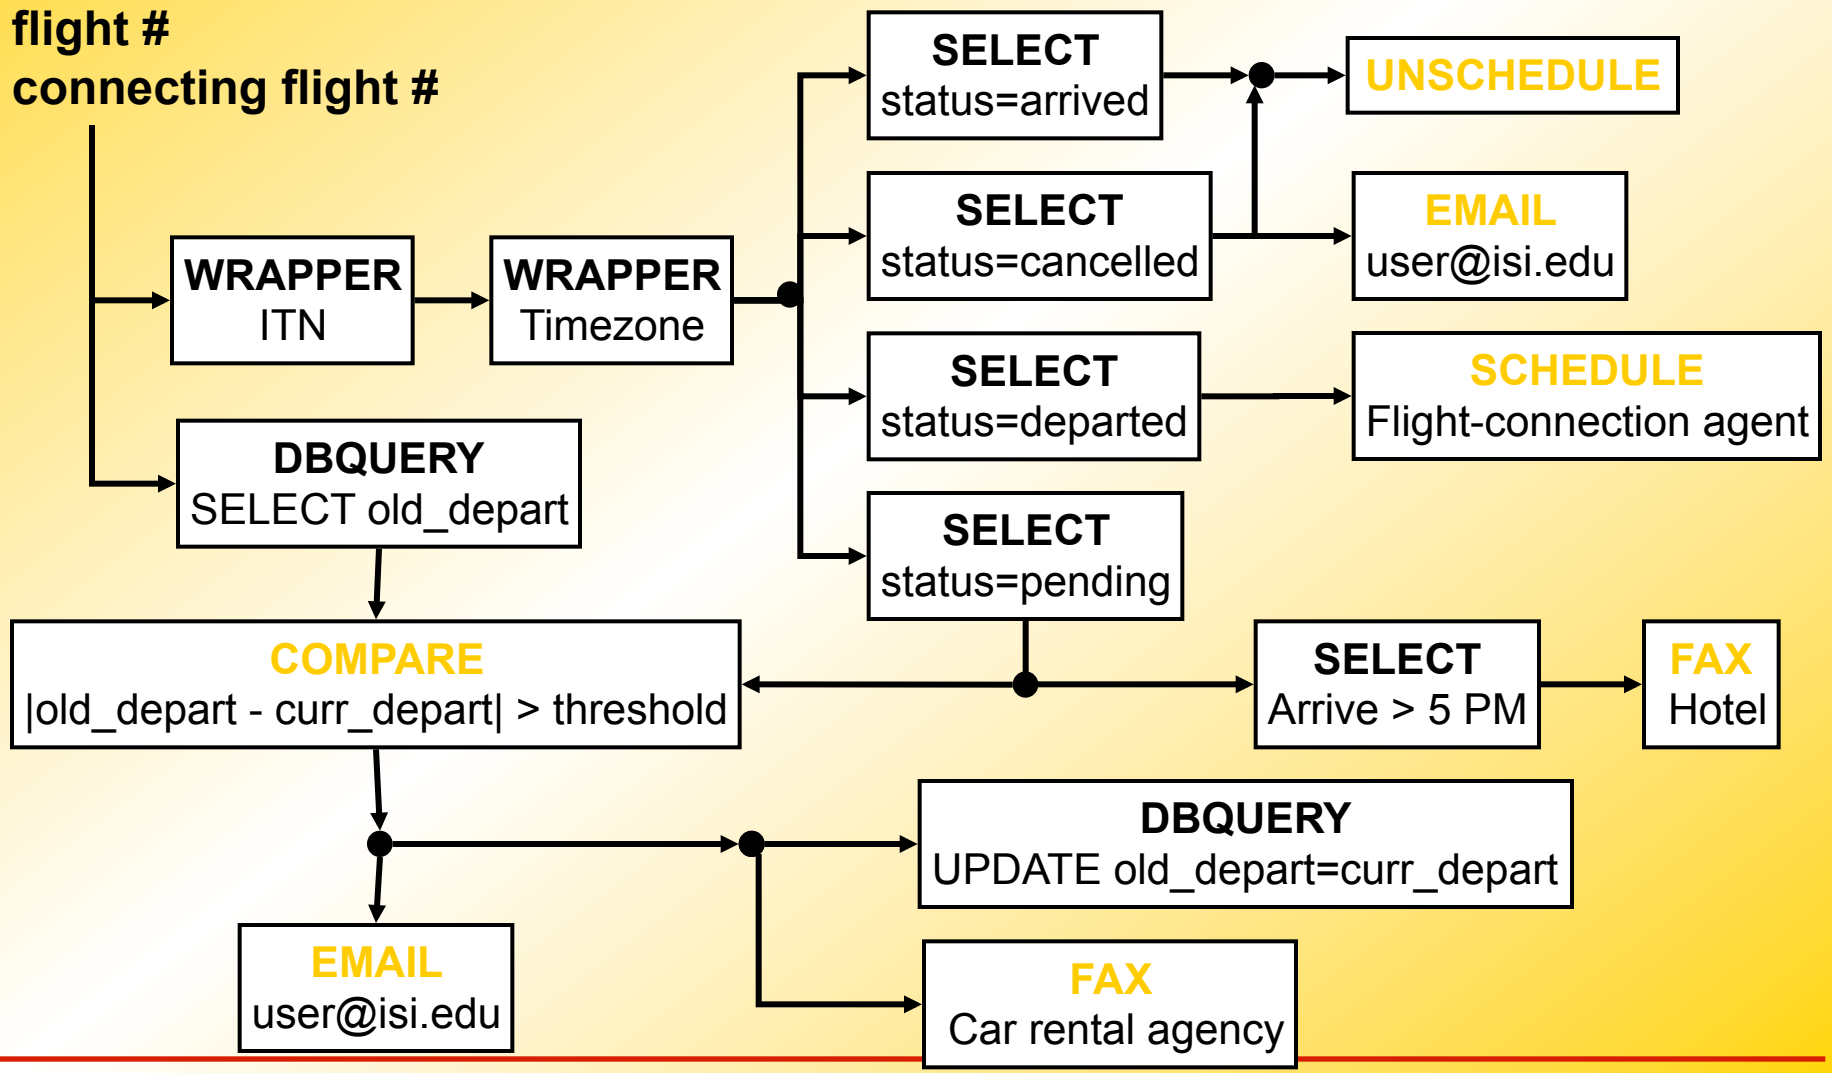

Outline

- The Electric Elves
  - Overview
- The Travel Elves
  - Constraint-based data integration
  - Efficient plan execution
- Beyond the Elves
  - Deployment to DARPA
  - Some lessons
    - Learning expectations
    - Finding alternative sources
    - Predicting prices
- Discussion

Deploying the Elves at DARPA

- The office elves were successful and received a lot of press
- DARPA decided they wanted a version running in their environment
- But there were issues
  - Integration into the DARPA environment
  - Privacy concerns
  - Software integration
  - …

The Travel Elves

- So we created the Travel Elves
  - No integration into the DARPA environment
  - Could run on ISI servers
  - Communicate via PDAs
- Trained the DARPA travel person to enter travel data using Heracles
- System deployed…
  - Used by program managers and office directors
  - …for a while

What Went Wrong and Why?

- Agent failures due to problems with the data
  - "Why didn't I receive any updates from the Elves?"
  - Source changed, unexpected data, source unavailable
- Agents failures due to problems in the logic
  - "Why did the message arrive after my flight left?"
  - Time zone exception
- Unexpected behaviors
  - "Why are the Elves pestering me with price changes?"
  - Frequent price changes
- Getting the definition "right"
  - "Those !*?&%* Elves sent a message to my cell phone at 3am to tell me my 8am flight was delayed!"
  - Agent worked as designed…
- Personalized agents
  - "Could you build an agent to do <X>?"
  - Every user wants something different

Outline

- The Electric Elves
  - Overview
- The Travel Elves
  - Constraint-based data integration
  - Efficient plan execution
- Beyond the Elves
  - Deployment to DARPA
  - Some lessons
    - Learning expectations
    - Finding alternative sources
    - Predicting prices
- Discussion

Learning Data Prototypes [Lerman et al., JAIR 2003]

- Approach to learning the structure of data
- Token-level syntactic description
  - descriptive but compact
  - computationally efficient
- Structure is described by a sequence (pattern) of general and specific tokens.
- Data prototype = starting & ending patterns

STREET_ADDRESS

220 Lincoln Blvd

420 S Fairview Ave

2040 Sawtelle Blvd

start with:

_NUM _CAPS

end with:

_CAPS Blvd

_CAPS _CAPS

Token Syntactic Hierarchy

- Tokens = words
- Syntactic types
  - e.g., NUMBER, ALPHA
- Hierarchy of types
  - allows generalization
- Extensible
  - new types
  - domain-specific information

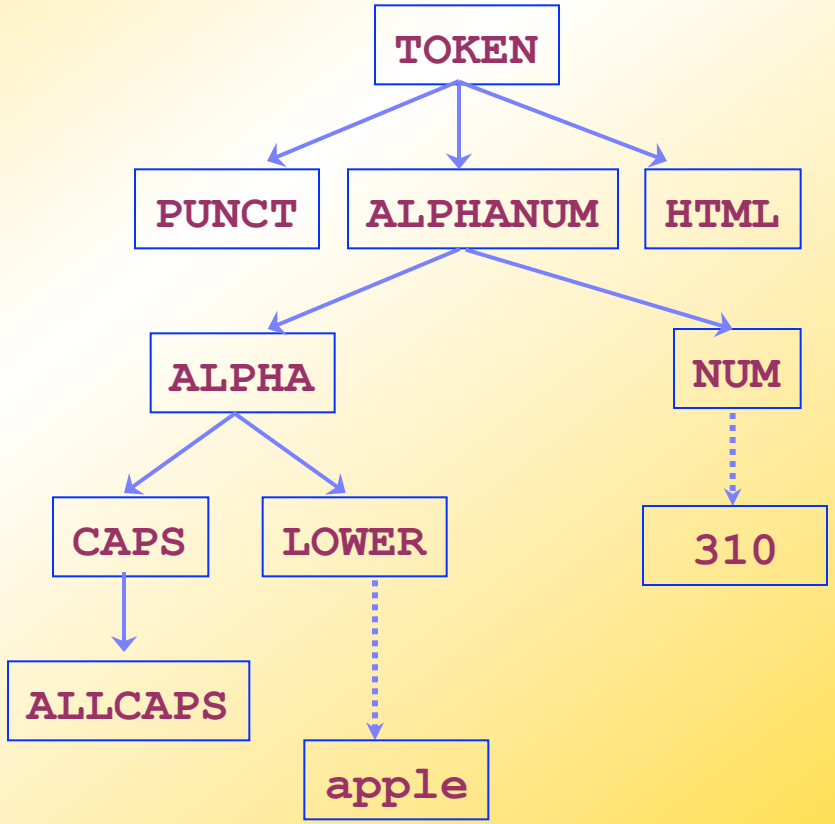

Prototype Learning Algorithm

- No explicit negative examples
- Learn from positive examples of data
- Find patterns that
  - describe many of the positive examples of data
  - highly unlikely to describe a random token sequence (implicit negative examples)
- are statistically significant patterns

  at α=0.05 significance level
- **DataPro** – efficient (greedy) algorithm

Outline

- The Electric Elves
  - Overview
- The Travel Elves
  - Constraint-based data integration
  - Efficient plan execution
- Beyond the Elves
  - Deployment to DARPA
  - Some lessons
    - Learning expectations
    - Finding alternative sources
    - Predicting prices
- Discussion

Learning New Sources of Data [Carman & Knoblock, 2005]

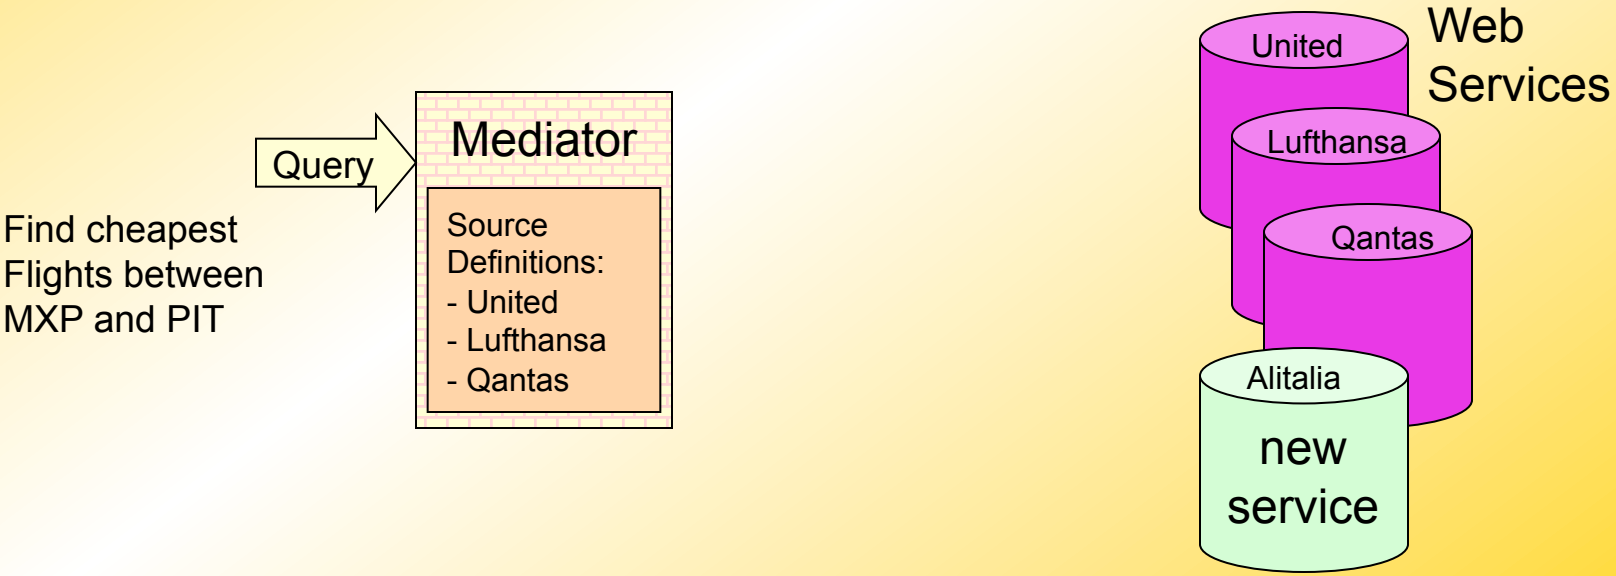

Learning New Sources of Data [Carman & Knoblock, 2005]

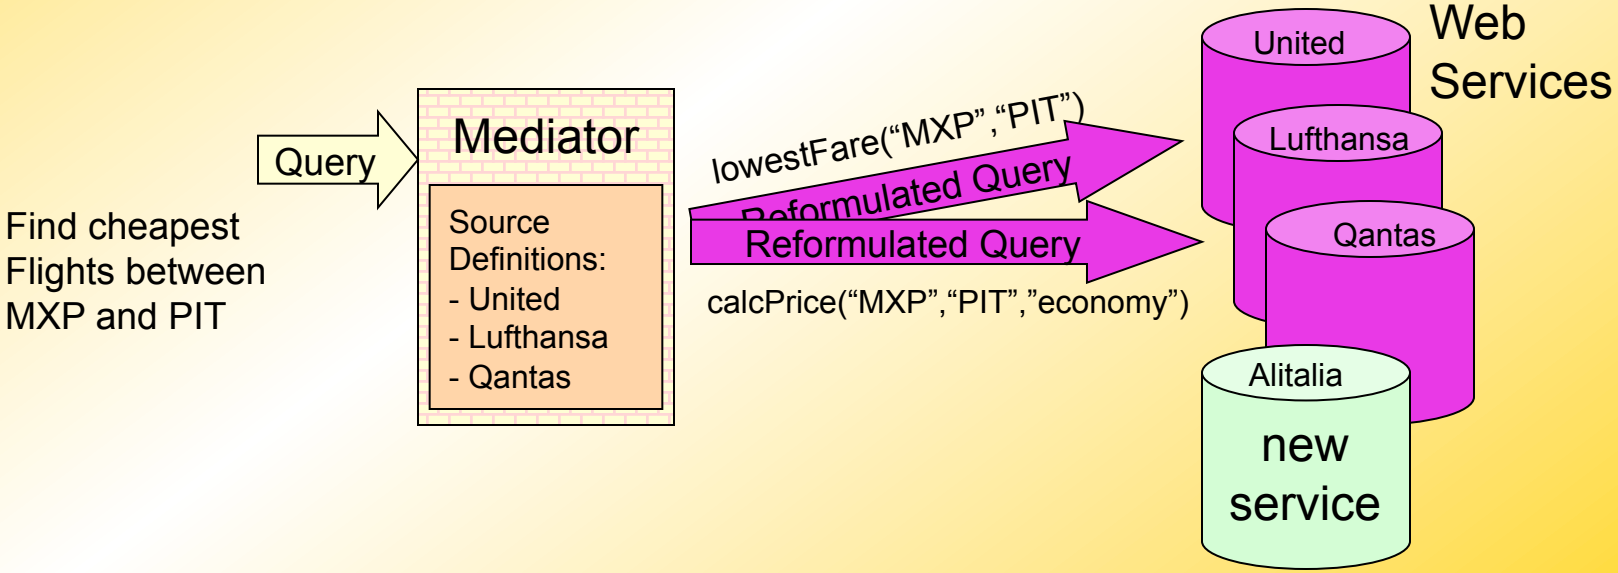

Learning New Sources of Data [Carman & Knoblock, 2005]

- Need source definitions to incorporate new data
- Time consuming and difficult to write these descriptions
- Can we discover definitions automatically?

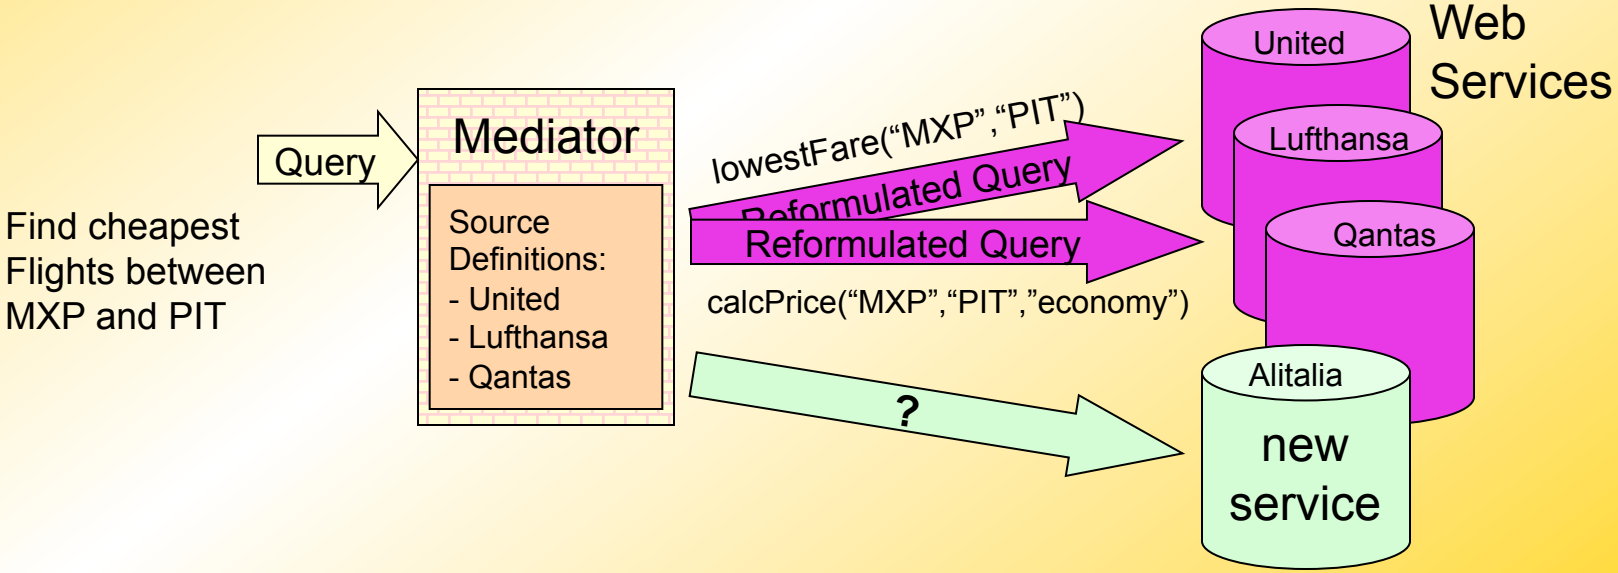

The Framework

Intuition: *Services often have similar semantics, so we should be able to use what we know to induce that which we don't*

Two phase algorithm

For each operation provided by the new service:

1. Classify its input/output data types
   - *Classify* inputs based on metadata similarity
   - *Invoke* operation & classify outputs based on data
2. Induce a source definition
   - *Generate* candidates via Inductive Logic Programming
   - *Test* individual candidates by reformulating them

Inducing Source Definitions: A Simple Example

Mediator

Semantic Types:
currency ⊆ {USD, EUR, AUD}
rate ⊆ {1936.2, 1.3058, 0.53177}

Predicates:
exchange(currency,currency,rate)

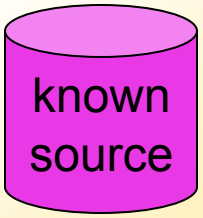

LatestRates($country1,$country2,rate):-
exchange(country1,country2,rate)

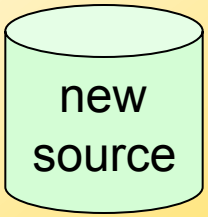

RateFinder($fromCountry,$toCountry,val):- ?

Inducing Source Definitions: A Simple Example

- Step 1: use metadata to classify input types

Mediator

Semantic Types:
currency ⊆ {USD, EUR, AUD}
rate ⊆ {1936.2, 1.3058, 0.53177}

Predicates:
exchange(currency,currency,rate)

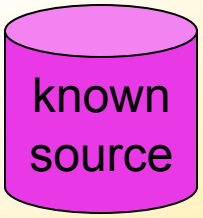

LatestRates($country1,$country2,rate):-
exchange(country1,country2,rate)

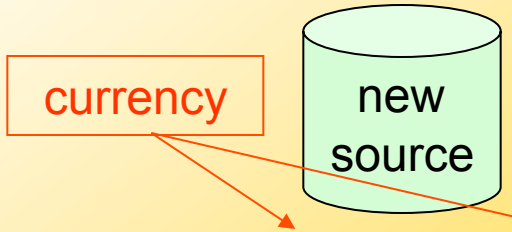

RateFinder($fromCountry,$toCountry,val):- ?

Inducing Source Definitions: A Simple Example

- Step 1: use metadata to classify input types
- Step 2: invoke service and classify output types

Semantic Types:
currency ⊆ {USD, EUR, AUD}
rate ⊆ {1936.2, 1.3058, 0.53177}

Predicates:
exchange(currency,currency,rate)

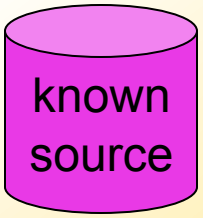

LatestRates($country1,$country2,rate):-
exchange(country1,country2,rate)

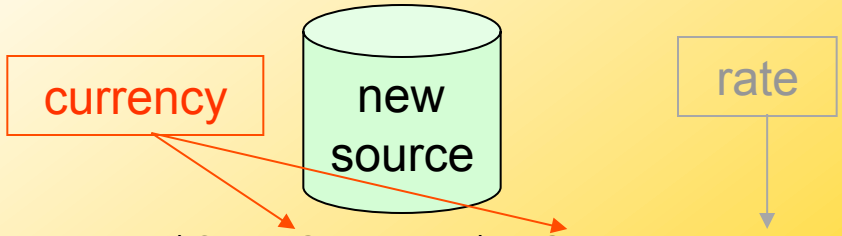

RateFinder($fromCountry,$toCountry,val):- ?

{<EUR,USD,1.30799>,<USD,EUR,0.764526>,…}

Inducing Source Definitions: A Simple Example

- Step 3: generate plausible source definitions

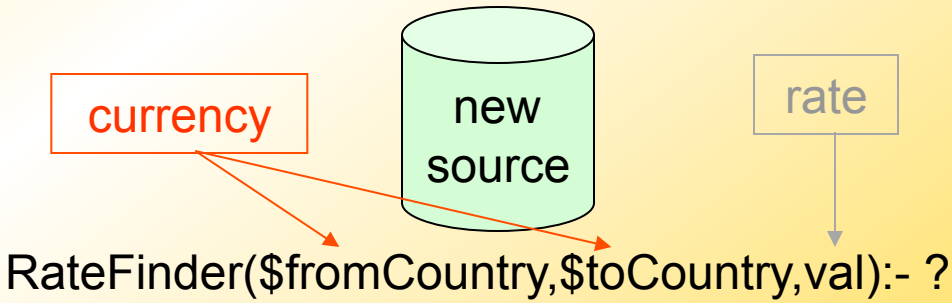

Predicates:

exchange(currency,currency,rate)

Inducing Source Definitions: A Simple Example

- Step 3: generate plausible source definitions

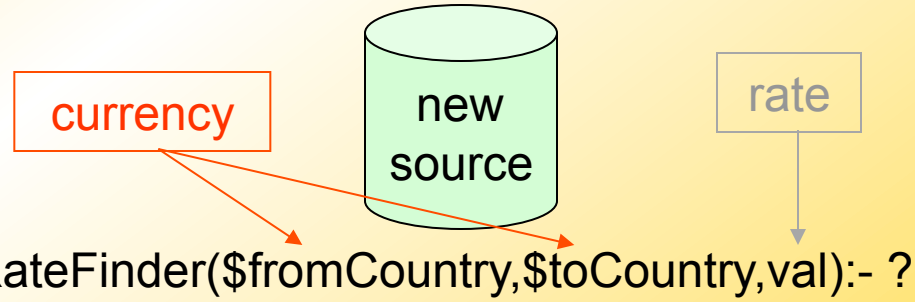

Mediator

def_1($from, $to, val) :- exchange(from,to,val)

def_2($from, $to, val) :- exchange(to,from,val)

Predicates:

exchange(currency,currency,rate)

Inducing Source Definitions: A Simple Example

- Step 3: generate plausible source definitions
- Step 4: reformulate in terms of other sources

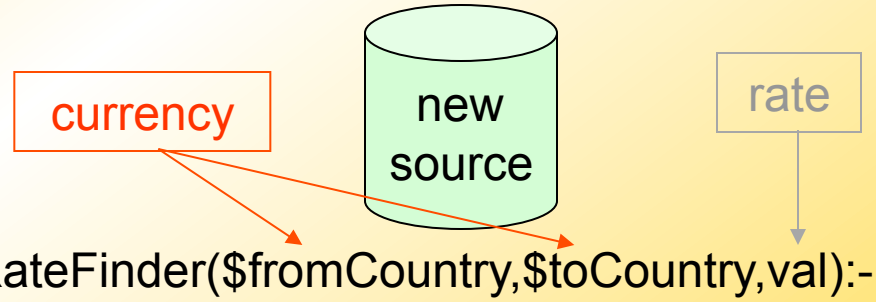

Mediator

Predicates:
exchange(currency,currency,rate)

def_1($from, $to, val) :- exchange(from,to,val)
def_2($from, $to, val) :- exchange(to,from,val)
def_1($from, $to, val) :- LatestRates(from,to,val)
def_2($from, $to, val) :- LatestRates(to,from,val)

Inducing Source Definitions: A Simple Example

- Step 3: generate plausible source definitions
- Step 4: reformulate in terms of other sources
- Step 5: invoke service and compare output

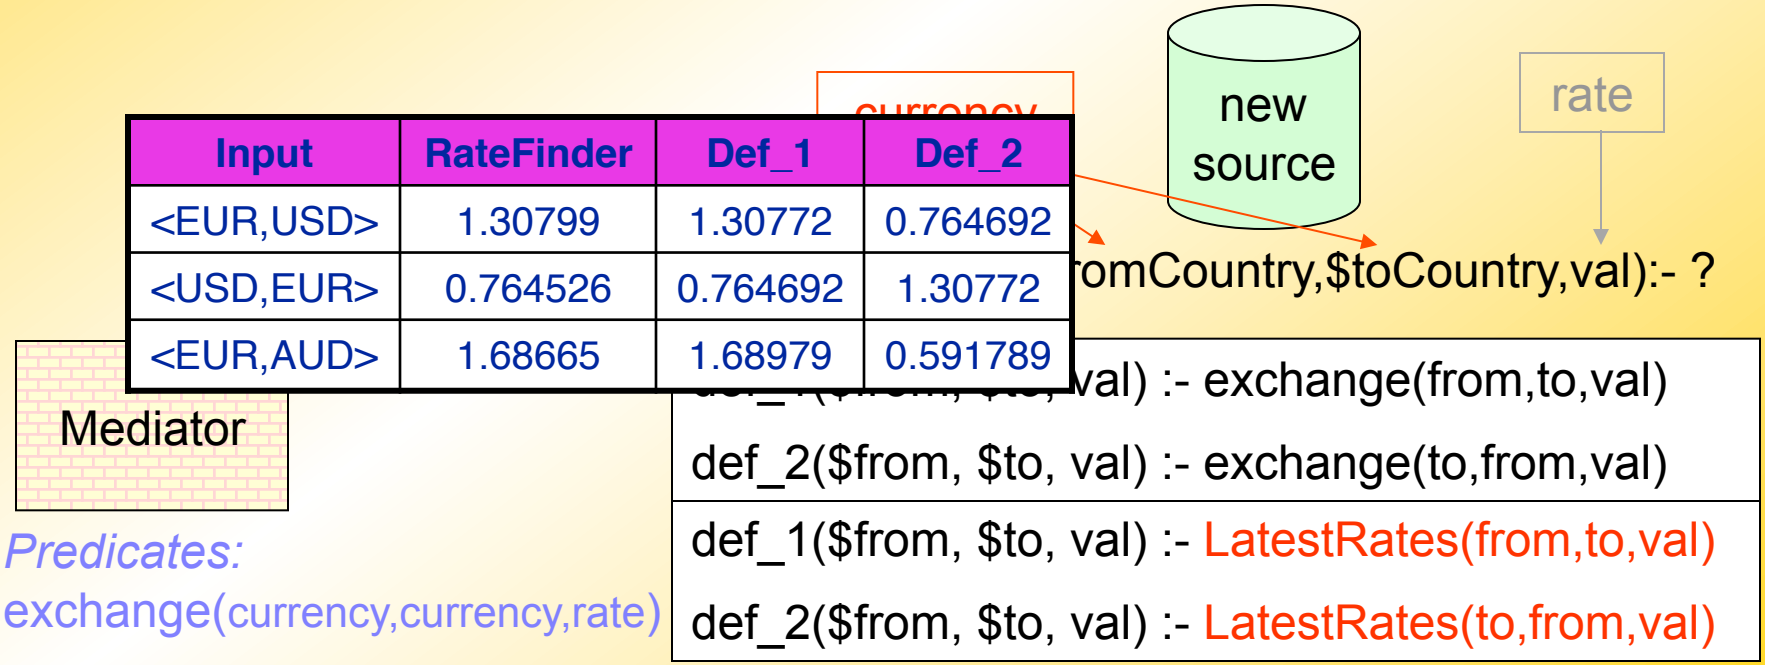

Input	RateFinder	Def_1	Def_2
<EUR,USD>	1.30799	1.30772	0.764692
<USD,EUR>	0.764526	0.764692	1.30772
<EUR,AUD>	1.68665	1.68979	0.591789

Inducing Source Definitions: A Simple Example

- Step 3: generate plausible source definitions
- Step 4: reformulate in terms of other sources
- Step 5: invoke service and compare output

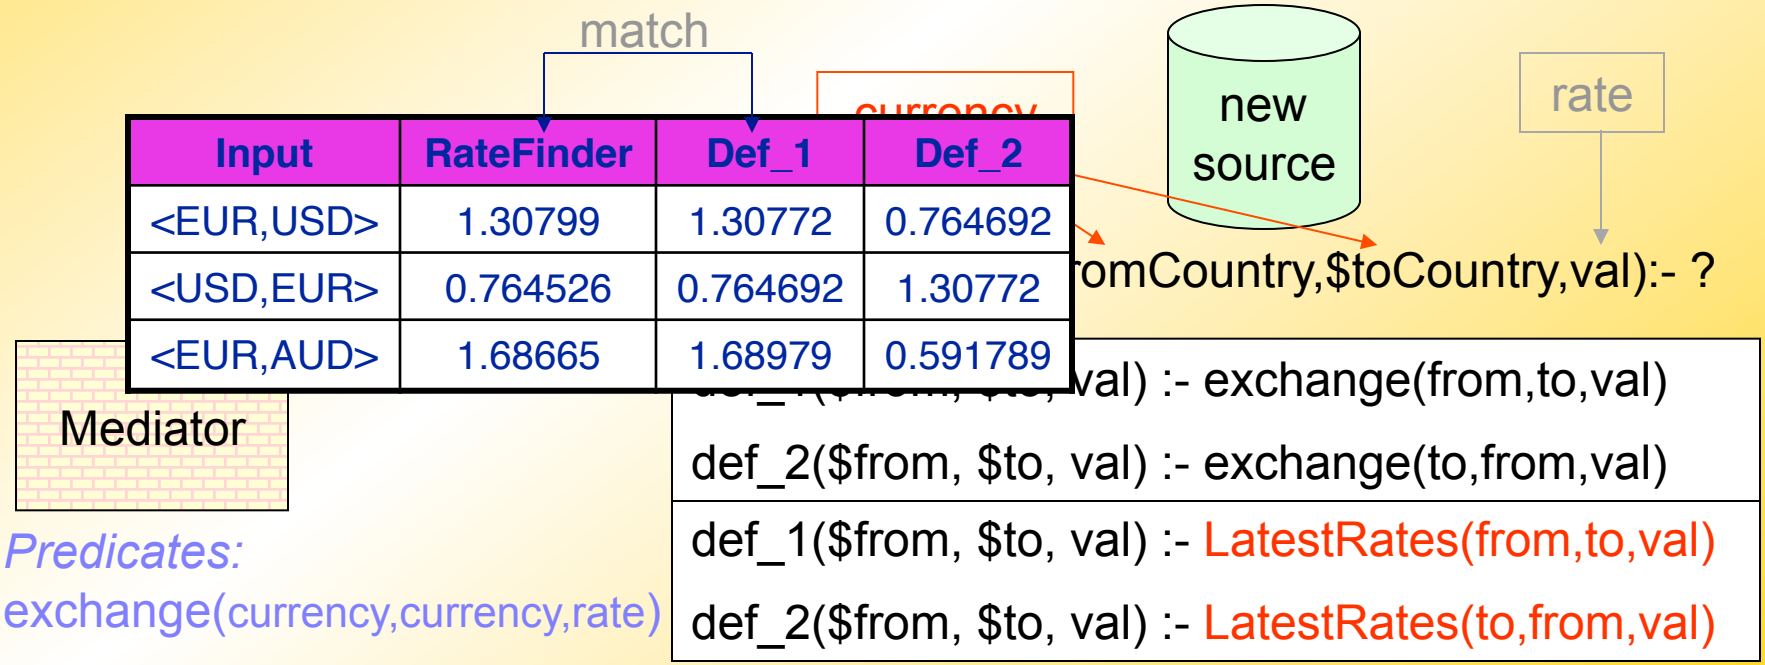

Outline

- The Electric Elves
  - Overview
- The Travel Elves
  - Constraint-based data integration
  - Efficient plan execution
- Beyond the Elves
  - Deployment to DARPA
  - Some lessons
    - Learning expectations
    - Finding alternative sources
    - Predicting prices
- Discussion

Predicting Airline Prices [Etzioni et al., KDD 2003]

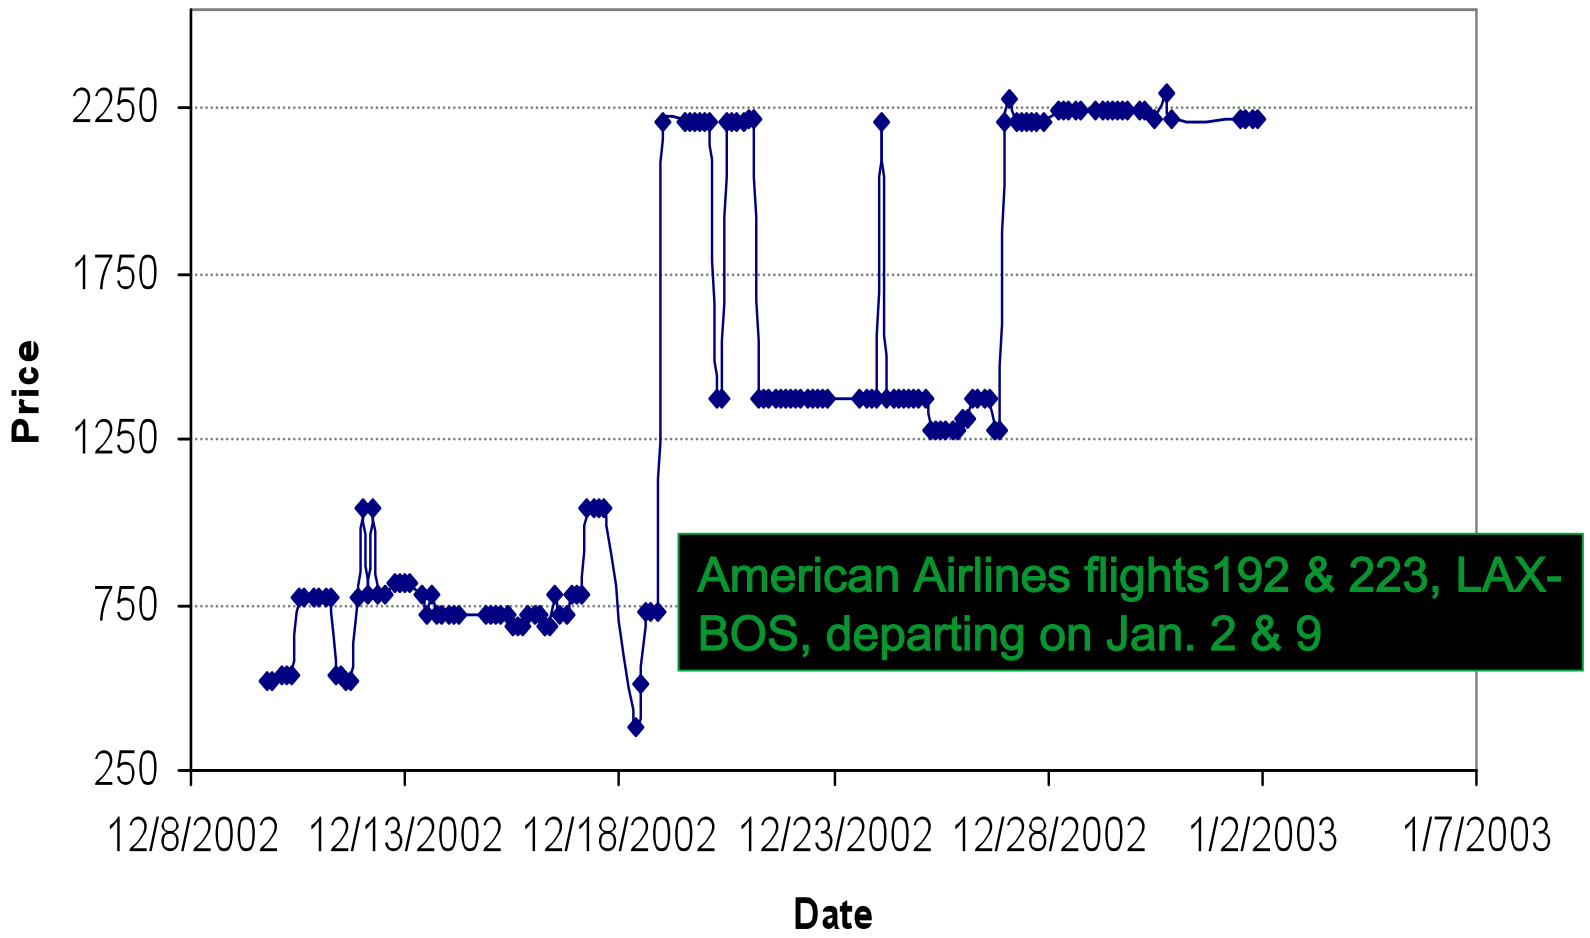

Hamlet: To Buy or Not to Buy

- Collected airline flight data over several months
- Developed a learning algorithm to predict whether to buy immediately or wait to buy a ticket
- Exploits the fact that airline pricing is done with a relatively static, but unknown algorithm
- Pricing can be learned by considering the pricing on the same flight on previous days

Data Set

- Extracted data from online sources using wrappers
- Collected over 12,000 price observations:
  - Lowest available fare for a one-week roundtrip
  - LAX-BOS and SEA-IAD
  - 6 airlines including American, United, etc.
  - 21 days before each flight, every 3 hours

Learning Algorithm

- Stacking with three base learners:
  1. Rule learning (Ripper) (e.g., R=wait)
  2. Time series
  3. Q-learning (e.g., Q=buy)
- Ripper used as the meta-level learner.
- Output: classifies each decision point as 'buy' or 'wait'.

Savings by Method

- **Net** savings = cost now – cost at purchase point.
- Penalty for sell out = upgrade cost. 0.42% of the time.
- Total ticket cost is $4,579,600.

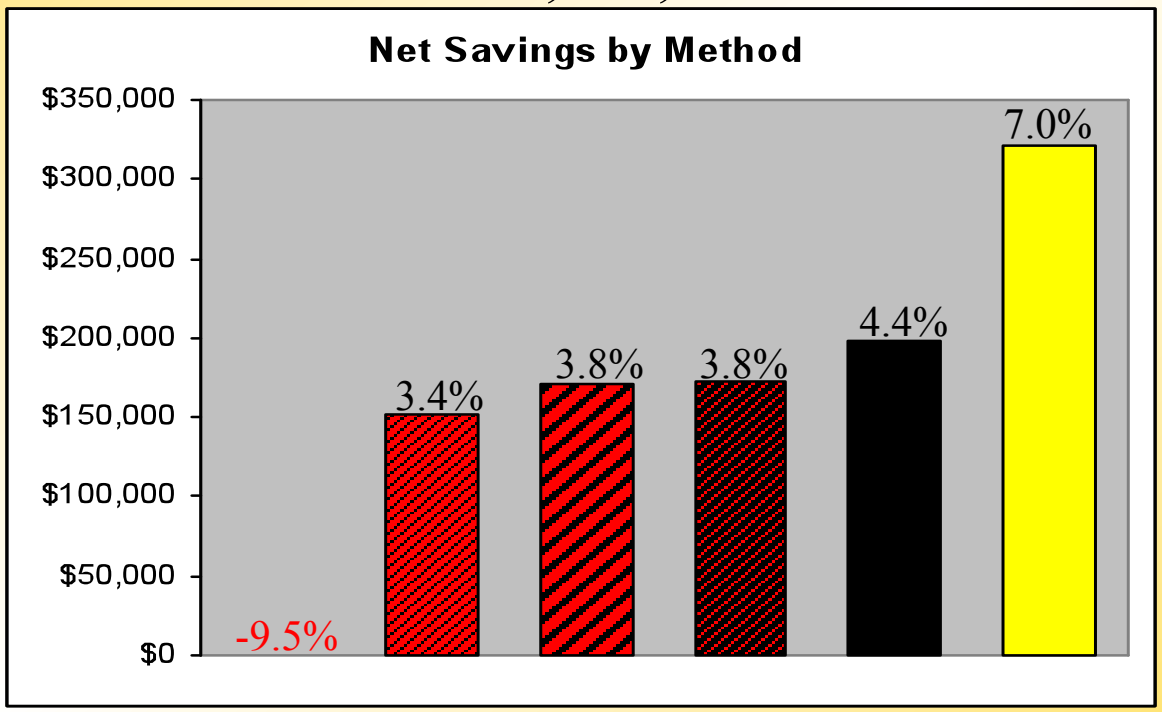

Status

- Patent granted on preditive pricing of air fares based on historical data
- Technology licensed to Hamlet, Inc. started by Oren Etzioni
- Company raised $7M in VC funding and is now called Farecast.
- Stay tuned…

Discussion

- We need to build agents that:
  - Robustly accomplish their tasks, responding appropriately to failures
  - Learn from their past experience
  - Rapidly build personalized agents without manual programming
  - Communicate flexibly with humans and software agents
  - Explain their behavior both on success or failure
  - Dynamically compose new agents and behaviors from existing agents