

MAYANK KEJRIWAL

4676 Admiralty Way, Ste. 1001, Marina Del Rey, CA 90292 | 1-217-819-6696 | kejriwal@isi.edu | kejriwalresearch.azurewebsites.net | <https://www.linkedin.com/in/mayankkejriwal/> | <https://github.com/mayankkejriwal>

CURRENT POSITION

Computer Scientist

Since June 2016

Information Sciences Institute
University of Southern California (USC) Viterbi School of Engineering

EDUCATION

University of Texas at Austin

M.Sc. and Ph.D. in Computer Science

Aug. 2012-May 2016

Current G.P.A: 3.83

Ph.D. thesis: *Populating a Linked Data Entity Name System*

Adviser: Daniel P. Miranker

Key Courses: Natural Language Processing, Information Retrieval, Neural Networks, Semantic Web, Algorithms, Graphical Models

University of Illinois at Urbana-Champaign

B.Sc. in Computer Engineering, Highest Honors

Aug. 2009-Aug. 2012

Graduation G.P.A: 3.89

Honors thesis: *Machine Learning Techniques in Offline Handwriting Transcription*

Advisers: P. S. Carney and Kenton McHenry

Upper-level Courses: Artificial Intelligence, Scientific Computation, Machine Learning, Computer Graphics, Theory of Computation, Computer Organization and Design, Digital Systems Laboratory, Analog Signal Processing, Solid State Electronic Devices, Computer Systems Engineering, Senior Research Project, Applied Linear Algebra, Probability with Engineering Applications, Differential Equations, Calculus III, Quantum Physics I, Classical Mechanics I and II, Electromagnetic Fields, Thermal and Statistical Physics, Classical Physics Lab

National University of Singapore

Study Abroad

Jan. 2012-May 2012

Key Courses: Introduction to Economics, Introduction to Japanese Studies, and Japan in the 20th Century

SCHOLARSHIPS, AWARDS AND TRAVEL GRANTS

AAAI Doctoral Consortium Travel Award	2015
Microsoft Azure for Research Grant	2014
National Science Foundation Travel Grant	2014
Department Nominee: Microsoft, Google fellowships	2013
Department Travel Grant	2013
MCD Fellowship	2012-2014
Daniel and Carol Dobberpuhl Award	2012
Senior 100 Honorary	2012
Henry O. Koehler Scholarship	2011
Chancellor's Scholar	2010-2012
Edmund J. James Scholar	2009-2012
Dean's List	2009-2012

EXPERIENCE

Information Sciences Institute, USC

Computer Scientist

June 2016-Current

Information Integration group

Capsenta, Inc.

Semantic Web Consultant

May 2016-June 2017

Worked on integrating advanced ontology matching solutions in the company's current products

University of Texas at Austin

Graduate Research Assistant

June 2014-May 2016

Research group of Daniel P. Miranker

CareerBuilder LLC, Norcross, Georgia

Data Scientist Intern

May 2015-Aug. 2015

Worked in the R&D team on the Recruitment Edge and Company Normalization products

University of Texas at Austin

Teaching Assistant

Aug. 2012-June 2014

Contemporary Issues in Computer Science (Spring 2014)

Automata Theory (Spring 2013)

Data Management (Falls 2013, 2012)

Artificial Intelligence (Fall 2012)

Rackspace, the Open Cloud Company, San Antonio, Texas

Summer Intern

May 2013-July 2013

Worked on the Big Data team helping with a new Hadoop-based product roll-out

National Center for Supercomputing Applications, Urbana, Illinois

Undergraduate Research Intern

May 2011-Aug. 2011

Developed scalable machine learning techniques to automate digital handwriting transcription of terabyte-level US census data

University of Illinois at Urbana-Champaign

Camp Counselor

May 2010-Aug. 2010

Designed STEM exercises for high-school students, and supervised them thereof

University of Illinois at Urbana-Champaign

Freshman Research Intern

May 2010-Aug. 2010

Helped to automate biomedical image processing tasks using the ImageJ software tool

SKILLS

- **Specialties** –Machine Learning, Data Integration, Semantic Web, Natural Language Processing
- **Languages**– Java, C/C++, Python, R, MATLAB, SQL, SPARQL, Scala, Latex
- **Systems/Tools**– Hadoop, Microsoft Azure, Eclipse IDE, Weka, LibSVM, SecondString, FEBRL, Lucene, Mallet

LANGUAGES

- **Native** – English, Hindi
- **Fluent** – Bengali

ACTIVITIES AND SERVICE

- **AI Seminar Coordinator:** Information Sciences Institute, 2017
- **Program Committee Member:** International Joint Conference on Artificial Intelligence (IJCAI), 2016.
- **National Science Foundation Data Science Workshop (2015):** an NSF-funded workshop attended by Ph.D. level data scientists from across the United States
- **Heidelberg Laureate Forum (2015):** Participant, one of among 200 globally selected Math and CS students
- **Professional Societies:** International Neural Network Society (INNS), Association for the Advancement of Artificial Intelligence (AAAI), Institute of Electrical and Electronics Engineers (IEEE), Association for Computing Machinery (ACM), Society for Industrial and Applied Mathematics (SIAM)
- **Honors Societies:** Phi Kappa Phi, Tau Beta Pi, Alpha Lambda Delta, Phi Eta Sigma
- **Volunteering and Hackathons:** Engineering Open House, 2010 (University of Illinois at Urbana-Champaign); dataHackUT, 2014 (University of Texas at Austin); Kaggle (2014-current); Internal hackathon, 2015 (CareerBuilder)
- **Reviewing/Sub-Reviewing:** ACM World Wide Web Conference (WWW), 2017; ACM Journal of Data and Information Quality (JDIQ); International Joint Conference on Artificial Intelligence (IJCAI), 2015; Very Large Databases (VLDB), 2013
- **Outreach:** AAAI Connections, February 2017 (AAAI, Iridescent Learning); *Present your research to a 12-year old*, 2015-2016 (University of Texas at Austin)

RESEARCH MENTOR

- **Haotian Zhang**
- **Rahul Kapoor**
- **Yixiang Yao**

BOOKS

Populating a Linked Data Entity Name System: A Big Data Solution for Unsupervised Instance Matching **2017**
Mayank Kejriwal
 IOS Press, Studies in the Semantic Web Series
 ISBN 978-3-89838-717-0 2017

Knowledge Graphs: Theory, Techniques and Applications **2018**
Mayank Kejriwal, Craig Knoblock, and Pedro Szekely
 MIT Press

PUBLICATIONS

Supervised Typing of Big Graphs using Semantic Embeddings **May 2017**
Mayank Kejriwal, Pedro Szekely
 Semantic Big Data (SBD) Workshop at ACM SIGMOD, 2017, held in Chicago, Illinois

Information Extraction in Illicit Domains **April 2017**
Mayank Kejriwal, Pedro Szekely
 ACM World Wide Web Conference, 2017, held in Perth, Australia

Local, Domain-independent Heuristics for the FEIII Challenge: Lessons and Observations **June 2016**
Mayank Kejriwal, Daniel P. Miranker
 DSMW Workshop at ACM SIGMOD, 2016, held in San Francisco, California

A Pipeline for Extracting and Deduplicating Domain-Specific Knowledge Bases **October 2015**
Mayank Kejriwal, Qiaoling Liu, Ferosh Jacob and Faizan Javed
 Industry track in the IEEE International Conference on Big Data, held in Santa Clara, California

Decision-making Bias in Instance Matching Model Selection **October 2015**
Mayank Kejriwal, Daniel P. Miranker
 The 14th International Semantic Web Conference, to be held in Bethlehem, Pennsylvania

An Unsupervised Instance Matcher for Schema-free RDF Data **July 2015**
Mayank Kejriwal, Daniel P. Miranker
 The Journal of Web Semantics

Sorted Neighborhood for Schema-free RDF Data **May 2015**
Mayank Kejriwal, Daniel P. Miranker
Winner of best paper award
 The 4th Knowledge Discovery and Data Mining meets Linked Open Data workshop,
 The 12th European Semantic Web Conference, held in Portoroz, Slovenia

Semi-supervised Instance Matching using Boosted Classifiers **May 2015**
Mayank Kejriwal, Daniel P. Miranker
 The 12th European Semantic Web Conference, held in Portoroz, Slovenia

Entity Resolution in a Big Data Framework **January 2015**
Mayank Kejriwal
 Doctoral consortium, the 29th Conference on Artificial Intelligence (AAAI), held in Austin, TX

Populating Entity Name Systems for Big Data Integration **October 2014**
Mayank Kejriwal
 The 13th International Semantic Web Conference, held in Riva Del Garda, Italy

A Two-Step Blocking Scheme Learner for Scalable Link Discovery **October 2014**
Mayank Kejriwal and Daniel P. Miranker
 The 9th International Workshop on Ontology Matching at the 13th International Semantic Web Conference, held in Riva Del Garda, Italy

Schema Matching over Relations, Attributes and Data Values **July 2014**
 Aibo Tian, **Mayank Kejriwal** and Daniel P. Miranker

The 26th International Conference on Scientific and Statistical Database Management, held in Aalborg, Denmark

An Unsupervised Algorithm for Learning Blocking Schemes

December 2013

Mayank Kejriwal and Daniel P. Miranker

The 13th IEEE International Conference on Data Mining, held in Dallas, Texas

Extended Scaled Neural Predictor for Improved Branch Prediction

August 2013

Zihao Zhou, **Mayank Kejriwal** and Risto Miikkulainen

The IEEE International Conference on Neural Networks, held in Dallas, Texas

A Framework to Access Handwritten Information within Large Digitized Paper Collections

October 2012

Liana Diesendruck, Luigi Marini, Rob Kooper, **Mayank Kejriwal** and Kenton McHenry

The 8th IEEE International Conference on eScience, held in Chicago, Illinois

ABSTRACTS, POSTERS AND PRESS

The Summer of Data Science

October 2015

Quoted in blog post by Kristin Tolle, Director of Data Science at Microsoft Research

Populating a Linked Data Entity Name System

August 2015

Mayank Kejriwal, Daniel P. Miranker

Poster at the Heidelberg Laureate Forum, held in Heidelberg, Germany

Unsupervised Instance Matching on Schema-free Linked Data

August 2015

Mayank Kejriwal, Daniel P. Miranker

Poster at the National Science Foundation Data Science Workshop, held in Seattle, Washington

Minimally Supervised Instance Matching: An Alternate Approach

May 2015

Mayank Kejriwal, Daniel P. Miranker

Poster and abstract at the 12th European Semantic Web Conference, held in Portoroz, Slovenia

The Math Behind...The Web of Linked Data

March 2015

Mayank Kejriwal

Poster for the SIAM: Math Matters, Apply It! program

A Sorted Neighborhood Workflow for the Semantic Web

January 2015

Mayank Kejriwal and Daniel P. Miranker

Poster and abstract at the 29th Conference on Artificial Intelligence (AAAI)

On Linking Heterogeneous Dataset Collections

October 2014

Mayank Kejriwal and Daniel P. Miranker

Poster and abstract at the 13th International Semantic Web Conference, held in Riva Del Garda, Italy

Populating Entity Name Systems for Big Data Integration

October 2014

Mayank Kejriwal and Daniel P. Miranker

Doctoral Consortium Poster at the 13th International Semantic Web Conference, held in Riva Del Garda, Italy

An Unsupervised Algorithm for Learning Blocking Schemes

December 2013

Mayank Kejriwal and Daniel P. Miranker

Poster at the 13th IEEE International Conference on Data Mining, held in Dallas, Texas

Extended Scaled Neural Predictor for Improved Branch Prediction

August 2013

Zihao Zhou, **Mayank Kejriwal** and Risto Miikkulainen

Poster at the IEEE International Conference on Neural Networks, held in Dallas, Texas

Digitization and Search: A Non-Traditional Use of HPC

October 2012

Liana Diesendruck, Luigi Marini, Rob Kooper, **Mayank Kejriwal** and Kenton McHenry

Poster and abstract at the 8th IEEE International Conference on eScience, held in Chicago, Illinois

Toward Free and Searchable Historical Census Images

September 2011

Kenton McHenry, Luigi Marini, **Mayank Kejriwal**, Rob Kooper and Peter Bajcsy

- Published as an article in the Electronic Imaging & Signal Processing section by the International Society for Optics and Photonics
- Published as a technical article in infoDOCKET

TECHNICAL REPORTS*On the Complexity of Sorted Neighborhood*

January 2015

Mayank Kejriwal and Daniel P. Miranker

Technical Report (arXiv:1501.01696)

A DNF Blocking Scheme Learner for Heterogeneous Datasets

January 2015

Mayank Kejriwal and Daniel P. Miranker

Technical Report (arXiv:1501.01694)

N-Way Heterogeneous Blocking

February 2014

Mayank Kejriwal and Daniel P. Miranker

Regular Technical Report (TR-14-06), University of Texas at Austin

TALKS*Populating a Linked Data Entity Name System*

April 2016

Ph.D. Final Defense, University of Texas at Austin

Committee: Daniel P. Miranker (chair), Raymond Mooney, Joydeep Ghosh, Risto Miikkulainen and Eric Price

Populating a Linked Data Entity Name System

April 2016

AI Seminar: Interview Talk

Information Sciences Institute, USC

Decision-making Bias in Instance Matching Model Selection

October 2015

The 14th International Semantic Web Conference, to be held in Bethlehem, Pennsylvania*Decision-making Bias in Instance Matching Model Selection*

October 2015

The 14th International Semantic Web Conference, to be held in Bethlehem, Pennsylvania*Sorted Neighborhood for Schema-free RDF Data*

May 2015

The 4th Knowledge Discovery and Data Mining meets Linked Open Data workshop,The 12th European Semantic Web Conference, held in Portoroz, Slovenia*Semi-supervised Instance Matching using Boosted Classifiers*

May 2015

The 12th European Semantic Web Conference, held in Portoroz, Slovenia*Entity Resolution in a Big Data Framework*

January 2015

Doctoral consortium, the 29th Conference on Artificial Intelligence (AAAI)*Populating a Linked Data-based Entity Name System*

November 2014

Ph.D. Oral Proposal, University of Texas at Austin

Committee: Daniel P. Miranker (chair), Raymond Mooney, Joydeep Ghosh, Risto Miikkulainen and Eric Price

Populating Entity Name Systems for Big Data Integration

October 2014

The 13th International Semantic Web Conference, held in Riva Del Garda, Italy*A Two-Step Blocking Scheme Learner for Scalable Link Discovery*

October 2014

The 9th International Workshop on Ontology Matching at the 13th International Semantic Web Conference, held in Riva Del Garda, Italy

Populating an Entity Name System

Research Preparation Exam, University of Texas at Austin
Committee: Daniel P. Miranker (chair), Greg Plaxton and Pradeep Ravikumar

March 2014

An Unsupervised Algorithm for Learning Blocking Schemes

The 13th IEEE International Conference on Data Mining, held in Dallas, Texas

December 2013

Machine Learning Techniques for Offline Handwriting Transcription

Undergraduate honors thesis presentation at the University of Illinois at Urbana-Champaign

December 2012